
Differentiable Sparsity via D -Gating: Simple and Versatile Structured Penalization

Chris Kolb^{1,2*}, Laetitia S. Frost¹, Bernd Bischl^{1,2}, and David Rügamer^{1,2}

¹Department of Statistics, LMU Munich

²Munich Center for Machine Learning (MCML)

Abstract

Structured sparsity regularization offers a principled way to compact neural networks, but its non-differentiability breaks compatibility with conventional stochastic gradient descent and requires either specialized optimizers or additional post-hoc pruning without formal guarantees. In this work, we propose D -Gating, a fully differentiable structured overparameterization that splits each group of weights into a primary weight vector and multiple scalar gating factors. We prove that any local minimum under D -Gating is also a local minimum using non-smooth structured $L_{2,2/D}$ penalization, and further show that the D -Gating objective converges at least exponentially fast to the $L_{2,2/D}$ -regularized loss in the gradient flow limit. Together, our results show that D -Gating is theoretically equivalent to solving the original group sparsity problem, yet induces distinct learning dynamics that evolve from a non-sparse regime into sparse optimization. We validate our theory across vision, language, and tabular tasks, where D -Gating consistently delivers strong performance–sparsity tradeoffs and outperforms both direct optimization of structured penalties and conventional pruning baselines.

1 Introduction

Sparsity in deep learning models has received considerable attention in recent years. On the one hand, *unstructured* sparsity methods remove individual weights to reduce parameter counts and achieve high compression ratios, but they produce irregular connectivity patterns that are hard to accelerate on standard hardware. On the other hand, *structured* sparsity targets entire groups of parameters, such as neurons [47, 60], convolutional filters [36, 38], or attention heads [11, 40, 51, 58, 68], yielding coarser sparsity patterns that translate directly into reductions in floating-point operations and memory requirements on existing hardware, which results in more efficient deployment of large models [5, 13].

Beyond computational advantages, introducing sparsity can also improve generalization performance [9] and increase model interpretability [19]. Nevertheless, most popular sparsification techniques in deep learning are not based on the non-smooth L_1 and $L_{2,1}$ penalties widely used in classical statistics and machine learning [53, 54], but rather constitute iterative pruning and retraining pipelines [2, 19, 32], whose main sparsification mechanism is defined by heuristic pruning criteria like parameter magnitudes. In these methods, the pruning step is decoupled from training, making it difficult to characterize precisely what overall objective is optimized and to provide principled guarantees. Further, the decision space is vast—pruning at initialization [9, 10, 17, 34, 52, 59], after training [36, 38, 39, 66], or sparsification during training [31, 44, 46], each with their own subtleties and tradeoffs [5, 13, 19]—making it cumbersome for practitioners to select a method that balances efficiency, accuracy, and theoretical soundness.

*Author correspondence to chris.kolb@stat.uni-muenchen.de




Figure 1: Parameter trajectories for a two-feature squared loss toy objective with non-convex $L_{2,2/D}$ regularization $\mathcal{L}(\mathbf{w}) = (y - x_1 w_1 - x_2 w_2)^2 + \lambda \|\mathbf{w}\|_2^{2/D}$ whose global minimizer is $(w_1^*, w_2^*) = (0, 0)$. **Left:** Failure of direct GD optimization to converge to $\mathbf{0}$ because of the non-differentiability at the origin. **Right:** D -gated objective where $\mathbf{w} = \boldsymbol{\omega} \cdot \prod_{d=1}^{D-1} \gamma_d$, converging smoothly to $\mathbf{0}$.

Sparsity penalties Structured sparsity penalties such as the $L_{2,1}$ norm are, in theory, capable of eliminating unimportant parameter groups, but in deep learning, they have mostly served as heuristics to steer post-hoc pruning rather than achieve exact sparsity [17, 38, 60]. Directly enforcing non-differentiable structured sparsity regularization requires solvers that can cope with its non-smooth nature; if this non-differentiability is ignored, optimization may oscillate or converge to dense, suboptimal solutions, as shown in Fig. 1. Replacing standard SGD with specialized procedures such as proximal-type algorithms (e.g. [7, 22, 41]) introduces substantial complexity, demands non-standard hyperparameter configurations, and often the routines are adapted to specific use cases or model classes, thereby foregoing modularity. This renders such approaches cumbersome to implement and inhibits their adoption for large-scale deep learning.

With these obstacles in mind, we ask and positively answer the following main question:

Research Question: Can we design a modular structured sparsity regularization method integrable into any architecture, amenable to SGD, with theoretical guarantees and little practical overhead?

1.1 Related literature

Structured sparsity A range of methods has been proposed to induce block-wise zeros in neural networks, yet they often rely on either post-hoc pruning or non-standard optimizers. Early work applies the convex $L_{2,1}$ penalty directly to weight groups but optimizes it with vanilla (sub-)gradient descent [47, 60], which fails to find sparse solutions and must be followed by an explicit pruning step [7]. To remedy the non-differentiability at zero, [7] proposes the use of a proximal algorithm, which we aim to avoid in favor of compatibility with SGD. [4] further generalizes the $L_{2,1}$ regularizer to non-convex $L_{2,q}$, $q < 1$, penalties using a custom optimizer. Rather than penalizing weight structures directly, [3, 22, 38] introduce shared scaling factors for each group and impose sparsity on those factors instead of the whole parameter group, but still require careful tuning. Although these competitors can yield exact sparsity under certain settings, they either fall back on pruning or abandon standard SGD, motivating our search for a fully differentiable, SGD-compatible alternative.

Differentiable sparsity A possible solution to incorporate sparsity-inducing penalties while retaining differentiability are approaches that split the parameters into multiple components and impose smooth L_2 regularization on the factors, which can be shown to induce the desired non-smooth sparsity penalty on the reconstructed parameters. This idea dates back to [14, 20] and has recently been adopted to incorporate differentiable L_1 -type sparsity regularization in neural networks [23, 28, 26, 27, 55, 69, 70]. In the case of matrix (instead of parameter) factorization using two factors, a low-rank bias, given by the trace norm, is induced on the product [25, 50]. The implicit bias literature also investigates such parameter decompositions *without* L_2 regularization, which can also induce sparsity under certain conditions—such as impractically small initialization scales [15, 56, 61, 67]. An extension to implicit group sparsity for linear models is presented by [37]. However, existing proposals either focus on unstructured sparsity, are constrained to only two factors, or do not constitute modular approaches applicable to arbitrary architectures. This leaves a gap in the




Figure 2: Overview of differentiable D -Gating method for structured sparsity (cf. Algorithm 1). For simplicity, we show D -Gating visually for a single fully-connected layer with input-wise grouping (colors), but our approach extends to arbitrary network components such as convolutional filters or attention heads. We proceed by applying D -Gating (red nodes and their connections) to the neural network weight and running SGD on the gating parameters with weight decay. After training, the weights are collapsed again and the zero structures removed, with the resulting sparse minimizers also being minimizers of the non-smooth $L_{2,2/D}$ -regularized objective.

current literature on whether extensions to arbitrary structures are possible, how such an approach can be implemented in practice, and to what extent there are theoretical guarantees to back this method.

1.2 Our contributions

Inspired by prior work on differentiable sparse regularization, we propose a new approach called D -Gating, which constitutes a structured sparsity-inducing penalty approach. It can be modularly incorporated in “arbitrary” architectures and neither incurs a notable overhead in additional parameters nor requires additional pruning steps, and is compatible with off-the-shelf SGD optimization. We further establish novel theoretical results that show the equivalence of our proposed differentiable regularization method and non-differentiable sparsity-inducing penalties (cf. Fig. 2), akin to what has been shown for approaches with unstructured sparsity penalties. Apart from theoretically and practically studying the loss landscape and training dynamics of our approach, we also validate our theory on an array of experiments to showcase its versatility in diverse deep learning applications.

2 Problem statement

In this paper, we propose a general structured sparsity approach for neural networks $f(\cdot, \mathbf{w})$ that allows penalizing excessive network components without placing restrictions on the type of architecture f or the position of the unit within the weight vector $\mathbf{w} \in \mathbb{R}^p$ that is targeted with the regularization. Structures such as filters naturally arise in neural networks, yielding a partition $\mathcal{G} = \{\mathcal{G}_1, \dots, \mathcal{G}_J\}$ of a subset of the indices $[p] := \{1, \dots, p\}$ of $\mathbf{w}$ into J groups $\mathbf{w}_{\mathcal{G}_j}$ with elements $w_{j,g}, g \in \mathcal{G}_j$. For filter sparsity in convolutional neural networks, $\mathcal{G}$ would be all indices for weights in the convolutional layers, and each $\mathcal{G}_j$ the indices of weights of one of the J filters.

Given this structure, we seek to optimize a general optimization problem

$$\underset{\mathbf{w} \in \mathbb{R}^p}{\text{minimize}} \quad \mathcal{L}_{\mathbf{w}}(\mathbf{w}) := \mathcal{L}_0(\mathbf{w}) + \lambda \|\mathbf{w}\|_{2,2/D}^{2/D} \quad (1)$$

for $D \geq 2$, where the unregularized objective $\mathcal{L}_0 = \sum_{i=1}^n \ell(y_i, f(\mathbf{x}_i, \mathbf{w}))$ is the sum of n observed loss contributions with loss $\ell: \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}^+$ evaluated on independent data points $\{\mathbf{x}_i, y_i\}_{i=1, \dots, n} \in (\mathcal{X} \times \mathcal{Y})^n$. The regularization term in Eq. (1) constitutes a generalization of what is often referred to as the group lasso [64], or $L_{2,1}$ penalty, which is recovered for $D = 2$. For $D > 2$, we obtain the more general and non-convex group penalty: $\|\mathbf{w}\|_{2,2/D}^{2/D} := \sum_{j=1}^J (\sum_{g \in \mathcal{G}_j} |w_{j,g}|^2)^{1/D}$ [21]. As this penalty is neither differentiable for $D = 2$ nor $D > 2$, SGD optimization of (1) yields unfavorable optimization dynamics and does not achieve exact sparsity. We therefore either require specialized optimization routines or a surrogate objective, which then induces the solution of (1). We choose the latter to create as little overhead (cf. Appendix E.2) and changes to training procedures as possible.

3 Differentiable structured sparsity via D -Gating

To solve Eq. (1) while enabling practitioners to use standard SGD optimizers, we derive a fully differentiable method that implicitly tackles Eq. (1) by employing an overparameterized model and a smooth surrogate penalty compatible with SGD optimization.

3.1 Model structure and gating variables

As we are interested in a general method for structured sparsity where arbitrary subsets of network weights can be sparsified, we assume a parametric learning model

$$f : \mathcal{X} \times \mathbb{R}^{q+p} \rightarrow \mathbb{R}^c, \quad (\mathbf{x}, \tilde{\mathbf{w}}) \mapsto f(\mathbf{x}; \tilde{\mathbf{w}}), \quad (2)$$

with inputs $\mathbf{x} \in \mathcal{X}$ and parameters $\tilde{\mathbf{w}} = (\mathbf{v}, \mathbf{w})$, where $\mathbf{w} \in \mathbb{R}^p$ contains the penalized parameters of interest and $\mathbf{v} \in \mathbb{R}^q$ the remaining parameters. Further, $\mathbf{w}$ is a partitioned (structured) weight object comprising J groups, $\mathbf{w} = (\mathbf{w}_j)_{j=1}^J \in \mathbb{R}^{p_1+\dots+p_J} = \mathbb{R}^p$.

To convert the non-smooth optimization problem into a smooth optimization problem, we subdivide $\mathbf{w}$ into two parts using the following gating operation:

Definition 1 (D -Gating). Let $\mathbf{w} \in \mathbb{R}^p$ and $\mathcal{G} = \{\mathcal{G}_1, \dots, \mathcal{G}_J\}$ be a partition of the indices $[p]$ of $\mathbf{w}$ into J groups. Further let $\boldsymbol{\omega} \in \mathbb{R}^p$ be the primary weight of the same size as $\mathbf{w}$, $\gamma_d \in \mathbb{R}$ one of $D-1$ vectors containing group-wise gating factors, and let $\gamma_j^\odot := \gamma_1 \odot \dots \odot \gamma_{D-1}$ denote the element-wise product of the gating factors with entries $\gamma_j^\odot = \prod_{d=1}^{D-1} \gamma_{j,d}$ for $j \in [J]$. For brevity, we collect the scaling factors in the matrix $\boldsymbol{\Gamma} = [\gamma_1, \dots, \gamma_{D-1}] \in \mathbb{R}^{J \times (D-1)}$. The D -Gating operation $\blacktriangleright : \mathbb{R}^p \times \mathbb{R}^{J \times (D-1)} \rightarrow \mathbb{R}^p, (\boldsymbol{\omega}, \boldsymbol{\Gamma}) \mapsto \boldsymbol{\omega} \blacktriangleright \gamma^\odot$, decomposes $\mathbf{w}$ as

$$\mathbf{w} = \boldsymbol{\omega} \blacktriangleright \gamma^\odot := \left(\boldsymbol{\omega}_j \prod_{d=1}^{D-1} \gamma_{j,d} \right)_{j=1}^J = \left(\boldsymbol{\omega}_j \cdot \gamma_j^\odot \right)_{j=1}^J, \quad (3)$$

and we call $\mathbf{w}$ D -Gated if it is parametrized as $\boldsymbol{\omega} \blacktriangleright \gamma^\odot$.

Intuitively, the D -Gating operation splits each group weight $\mathbf{w}_j$ into D factors: the vector $\boldsymbol{\omega}_j$, corresponding to the original group weights, and the $D-1$ additional gating factors $\gamma_{j,d} \in \mathbb{R}$, which are applied multiplicatively to all entries of $\boldsymbol{\omega}_j$.

3.2 Differentiable penalty

Given the gated formulation, we can now impose surrogate L_2 regularization on $(\boldsymbol{\omega}, \boldsymbol{\Gamma})$, defined as

$$\mathcal{L}(\mathbf{v}, \boldsymbol{\omega}, \boldsymbol{\Gamma}) = \mathcal{L}_0(\mathbf{v}, \boldsymbol{\omega} \blacktriangleright \gamma^\odot) + \lambda \mathcal{R}(\boldsymbol{\omega}, \boldsymbol{\Gamma}) \quad (4)$$

$$= \sum_{i=1}^n \ell(y_i, f(\mathbf{x}_i, (\mathbf{v}, \boldsymbol{\omega} \blacktriangleright \gamma^\odot))) + \frac{\lambda}{D} \underbrace{\sum_{j=1}^J (\|\boldsymbol{\omega}_j\|_2^2 + \sum_{d=1}^{D-1} \gamma_{j,d}^2)}_{\|\boldsymbol{\omega}\|_2^2 + \|\boldsymbol{\Gamma}\|_F^2} \quad (5)$$

where $\mathcal{L}_0(\mathbf{v}, \mathbf{w})$ denotes the unregularized, differentiable loss function with per-sample loss ℓ . In the following, we will denote local minimizers of the D -Gating objective as

$$(\hat{\mathbf{v}}, \hat{\boldsymbol{\omega}}, \hat{\boldsymbol{\Gamma}}) \in \arg \min_{(\mathbf{v}, \boldsymbol{\omega}, \boldsymbol{\Gamma}) \in \mathbb{R}^{q+p+J(D-1)}} \mathcal{L}(\mathbf{v}, \boldsymbol{\omega}, \boldsymbol{\Gamma}). \quad (6)$$

While the function values of Eq. (4) are not necessarily equal to those of Eq. (1), our next section provides a theoretical guarantee of the equivalence of both objectives with regards to their minima and shows that the difference between both objectives is vanishing at least exponentially fast.

4 Theoretical results

Because the presence of ungated parameters $\mathbf{v}$ is inconsequential for our analyses and all results directly carry over, they will be omitted from now on, and we assume for simplicity of exposition that all model parameters are D -gated.

4.1 Stationarity condition and loss simplification

The following result establishes that all stationary points of $\mathcal{L}(\omega, \Gamma)$ correspond to balanced D -gating parameters. Otherwise, one could continuously perturb the D -gating parameters toward a more balanced configuration without altering the effective parameter $\mathbf{w} = \omega \blacktriangleright \gamma^\odot$ while strictly decreasing the L_2 penalty $\mathcal{R}(\omega, \Gamma)$.

Lemma 1 (Balancedness at stationary points). *Let (ω, Γ) be D -gating parameters satisfying $\mathbf{w}_j = \omega_j \prod_{d=1}^{D-1} \gamma_{j,d}$ for $j \in [J]$. If (ω, Γ) is a stationary point of the L_2 -regularized objective $\mathcal{L}(\omega, \Gamma)$, then for each group j the gating factors are balanced in the sense that*

$$\|\omega_j\|_2^2 = \gamma_{j,1}^2 = \dots = \gamma_{j,D-1}^2 = \|\mathbf{w}_j\|_2^{2/D}. \quad (7)$$

Notably, the loss evaluated at balanced parameters simplifies to reveal its sparsity-inducing nature:

Corollary 1 (Loss simplification at balanced gating parameters). *Let (ω, Γ) be balanced D -gating parameters satisfying $\mathbf{w}_j = \omega_j \prod_{d=1}^{D-1} \gamma_{j,d}$ and Eq. (7) for $j \in [J]$. The D -gated objective $\mathcal{L}(\omega, \Gamma)$ in Eq. (5) then simplifies to*

$$\mathcal{L}_0(\omega \blacktriangleright \gamma^\odot) + \frac{\lambda}{D} (\|\omega\|_2^2 + \|\Gamma\|_F^2) = \mathcal{L}_0(\mathbf{w}) + \lambda \sum_{j=1}^J \|\mathbf{w}_j\|_2^{2/D} = \mathcal{L}_0(\mathbf{w}) + \lambda \|\mathbf{w}\|_{2,2/D}^{2/D} =: \mathcal{L}_{\mathbf{w}}(\mathbf{w}) \quad (8)$$

4.2 Equivalence of optimization problems

The previous result is reassuring as it implies the equivalence of objectives at balanced gating parameters. It does, however, not guarantee that optimizing one objective provides a meaningful solution for the other objective. Fortunately, such an equivalence can be established.

Theorem 1 (Equivalence of D -gating and $L_{2,2/D}$ regularization). *The two optimization problems*

$$\underset{\omega \in \mathbb{R}^p, \Gamma \in \mathbb{R}^{J \times (D-1)}}{\text{minimize}} \quad \mathcal{L}(\omega, \Gamma) := \mathcal{L}_0(\omega \blacktriangleright \gamma^\odot) + \frac{\lambda}{D} (\|\omega\|_2^2 + \|\Gamma\|_F^2) \quad (9)$$

$$\underset{\mathbf{w} \in \mathbb{R}^p}{\text{minimize}} \quad \mathcal{L}_{\mathbf{w}}(\mathbf{w}) := \mathcal{L}_0(\mathbf{w}) + \lambda \|\mathbf{w}\|_{2,2/D}^{2/D} \quad (10)$$

are equivalent in the sense that their local minima are identical. If $(\hat{\omega}, \hat{\Gamma}) \in \arg \min \mathcal{L}(\omega, \Gamma)$, then $\hat{\omega} \blacktriangleright \hat{\gamma}^\odot = \hat{\mathbf{w}} \in \arg \min \mathcal{L}_{\mathbf{w}}(\mathbf{w})$, and likewise, if $\hat{\mathbf{w}} \in \arg \min \mathcal{L}_{\mathbf{w}}(\mathbf{w})$, then any balanced gating representation $(\hat{\omega}, \hat{\Gamma})$ such that $\hat{\mathbf{w}} = \hat{\omega} \blacktriangleright \hat{\gamma}^\odot$ is a local minimizer of $\mathcal{L}(\omega, \Gamma)$.

Specifically, there is a bijective mapping between the local minimizers of $\mathcal{L}_{\mathbf{w}}(\mathbf{w})$ and the equivalence class of local minimizers of $\mathcal{L}(\omega, \Gamma)$.

4.3 Optimization dynamics

Under ubiquitous (S)GD-based optimization of the D -Gated objective in Eq. (5), we can additionally establish results characterizing the evolution of parameter balancedness, i.e., quantifying how fast the D -Gated objective converges to the original $L_{2,2/D}$ regularized loss.

4.3.1 Gradient flow dynamics of imbalances for D -gated models

The group-wise gradient flow dynamics for $j \in [J]$ are given by

$$\dot{\omega}_j = -\nabla_{\omega_j} \mathcal{L}, \quad \dot{\gamma}_{j,d} = -\partial_{\gamma_{j,d}} \mathcal{L}. \quad (11)$$

The gradients with respect to the D -Gating parameters ω_j and the $\gamma_{j,d}$ become, using the chain rule,

$$\nabla_{\omega_j} \mathcal{L} = \gamma_j^\odot \nabla_{\mathbf{w}_j} \mathcal{L}_0 + \frac{2\lambda}{D} \omega_j, \quad \partial_{\gamma_{j,d}} \mathcal{L} = \left(\frac{\gamma_j^\odot}{\gamma_{j,d}} \right) \omega_j^\top \nabla_{\mathbf{w}_j} \mathcal{L}_0 + \frac{2\lambda}{D} \gamma_{j,d} \quad d \in [D-1]. \quad (12)$$

We define the pair-wise imbalance $\mathcal{I}$ between any two group-wise factors $d \neq d' \in [D]$ and show it vanishes exponentially in time:

$$\mathcal{I}_{j,d,d'}(t) := \begin{cases} \|\omega_j(t)\|_2^2 - \gamma_{j,d'}(t)^2, & \text{if } d = 1, \\ \gamma_{j,d}(t)^2 - \gamma_{j,d'}(t)^2, & \text{if } d \neq 1. \end{cases} \quad (13)$$




Figure 3: Evolution of imbalance during SGD of a neuron-wise D -gated LeNet-300-100 for $D \in \{2, 3, 4\}$ (left to right). As predicted by our theory, the losses converge exponentially, with the rate increasing in λ and decreasing in D .

Lemma 2 (Exponential decay of imbalance). *Under the gradient flow dynamics Eqs. (11) and (12), the pair-wise imbalance $\mathcal{I}_{j,d,d'}(t)$ (13) between two gating parameters d, d' of group j satisfies*

$$\frac{d}{dt} \mathcal{I}_{j,d,d'}(t) = -\frac{4\lambda}{D} \mathcal{I}_{j,d,d'}(t) \quad \forall d \neq d', j \in [J]. \quad (14)$$

Solving the ODE shows $\mathcal{I}_{j,d,d'}(t)$ decays exponentially for $\lambda > 0$: $\mathcal{I}_{j,d,d'}(t) = \mathcal{I}_{j,d,d'}(0)e^{-\frac{4\lambda}{D}t}$.

For $\lambda = 0$, $\mathcal{I}_{j,d,d'}(t)$ is a conserved quantity [30]. On the loss or penalty level, the overall *misalignment* between both losses can be defined as $\mathcal{M}(\omega, \Gamma) := D^{-1}(\|\omega\|_2^2 + \|\Gamma\|_F^2) - \|\omega \blacktriangleright \gamma^\odot\|_{2,2/D}^{2/D} \geq 0$. Using this, the previous result can be extended to show $|\mathcal{L}(\omega, \Gamma) - \mathcal{L}_w(\omega \blacktriangleright \gamma^\odot)| \rightarrow 0$:

Lemma 3 (Convergence of D -gated loss to $L_{2,2/D}$ regularized loss). *Assume that the model parameters of Eq. (4) with effective weight $\mathbf{w}(t) = \omega(t) \blacktriangleright \gamma^\odot(t)$ as in Eq. (3) evolve with time t according to the gradient flow in Eqs. (11) and (12). Then, the D -Gated loss $\mathcal{L}(\omega(t), \Gamma(t))$ in Eq. (9) converges to the non-smooth $L_{2,2/D}$ regularized loss $\mathcal{L}_w(\mathbf{w}(t))$ in Eq. (10), evaluated at the effective weight $\mathbf{w}(t) = \omega(t) \blacktriangleright \gamma^\odot(t)$, at least exponentially fast given an initialization-dependent constant C :*

$$\mathcal{L}(\omega(t), \Gamma(t)) - \mathcal{L}_w(\mathbf{w}(t)) \leq Ce^{-\frac{4\lambda}{D}t}, \quad (15)$$

Intuitively, this is because balancedness in D -Gating is precisely the condition required for $\mathcal{L}(\omega, \Gamma)$ to simplify to $\mathcal{L}_w(\mathbf{w})$ (cf. Corollary 1). Hence, as the pair-wise imbalances vanish, the balancedness condition becomes increasingly true, and the two losses converge. For discrete-time imbalance dynamics, the dynamics becomes more convoluted, but we can establish geometric decay up to first-order in η , and find symmetry-induced absorbing SGD states [69] for balanced gating configurations.

Lemma 4 (Imbalance dynamics under SGD). *Consider the D -Gated objective Eq. (5). Then, under (S)GD, for any $j \in [J]$, (i) the pair-wise imbalances in Eq. (13) evolve as*

$$\mathcal{I}_{j,d,d'}^{(t+1)} = (1 - 4\lambda\eta/D) \mathcal{I}_{j,d,d'}^{(t)} + \eta^2 \Delta_{j,d,d'}^{(t)}, \quad \eta > 0, \quad d, d' \in [D], d \neq d', \quad (16)$$

with separate second-order terms $\Delta_{j,d,d'}$ for $d = 1$ and $d, d' > 1$. For sufficiently small η , $\mathcal{I}_{j,d,d'}^{(t+1)} \approx (1 - 4\lambda\eta/D) \cdot \mathcal{I}_{j,d,d'}^{(t)}$ exhibits discrete exponential decay. (ii) Balancedness is conserved between any two scalar factors $d, d' > 1$, i.e., $\mathcal{I}_{j,d,d'}^{(t)} = 0 \Rightarrow \mathcal{I}_{j,d,d'}^{(t')} = 0 \forall t' > t$, and (iii), for balanced zero representations $(\omega_j^{(t)}, \{\gamma_{j,d}^{(t)}\}_{d=1}^{D-1}) = \mathbf{0}$, it holds $(\omega_j^{(t')}, \{\gamma_{j,d}^{(t')}\}_{d=1}^{D-1}) = \mathbf{0} \forall t' > t$.

5 Numerical experiments

In the following, we empirically investigate the learning dynamics of our approach in Section 5.1 and then showcase various applications in Section 5.2 to demonstrate our method's modularity.

5.1 Learning dynamics and misalignment

5.1.1 Exponential decay of imbalance

We first validate our theoretical results on learning dynamics and loss convergence of D -Gating from Section 4.3. For this, we apply D -Gating with $D \in \{2, 3, 4\}$ to a LeNet-300-100 at the neuron level and train the model on the MNIST dataset using SGD. We use a grid of λ values and measure the loss convergence as defined in Eq. (15). Fig. 3 visualizes the results and confirms our theoretical findings on the exponential decay of the loss difference.




Figure 4: Regularization paths for sparse linear regression task using D -Gating. **Left:** Test RMSE vs λ . The curves for 2-Gating and group lasso coincide beyond a certain λ , but are outperformed by D -Gating with $D > 2$. **Middle:** Group sparsity of 2-Gating coincides with group lasso solution. Dashed grey line indicates optimal λ for all models. Deeper gating yields sparser solutions. **Right:** Transition of 2-Gating to group lasso solution beyond a certain λ coincides with zero imbalance attained after training. Direct optimization with GD yields notably different regularization paths far from the global minima. The dashed black line indicates 0 misalignment at the end of training. Means and 95% confidence intervals over ten simulations are shown for the left two plots.

5.1.2 D-gating and misalignment for group lasso

To further validate the equivalence of optimization problems as established in the previous section, we run a sparse linear regression where direct $L_{2,1}$ -regularized optimization is more accessible due to the availability of specialized optimization routines. For this, we simulate data as described in Appendix C.1 with 40 feature groups of 5 features each, of which 7 are informative (with truly non-zero effects). We use accelerated proximal gradient descent [48] to directly optimize the original $L_{2,1}$ -penalized linear model and apply our approach for $D \in \{2, 3, 4\}$ over the same grid of λ values. In addition, we also perform direct GD optimization of the $L_{2,1}$ -regularized linear model and compare all methods against an oracle (a linear model using only the signal variables). Results in Fig. 4 confirm the established equivalence between the original and D -gated objective for $D = 2$, but also demonstrate the improvement in the performance-sparsity tradeoff for $D > 2$.

5.2 Modularity

Next, we demonstrate the flexibility of our method. To this end, we study different types of structured sparsity problems arising in neural networks. In these experiments, we focus on demonstrating the broad applicability of our method rather than an exhaustive benchmark comparison. Our method supports any form of structured sparsity in neural networks, enabling diverse applications. Selected use cases are shown below. Further results are presented in Appendix E.

5.2.1 Feature selection in non-linear models

We start by investigating input feature selection, i.e., by individually gating the first-layer weights outgoing from each input feature. We follow the setup of [35] by running the proposed method,




Figure 5: Comparison of feature selection methods. Means and std. over 5 random initializations are reported.

LassoNet, as well as *HSIC* [63] on six diverse datasets as done in [35]. We use the same LeNet-300-100 architecture as a backbone for *LassoNet*, *HSIC* and our *D*-Gating approach (cf. Appendix C.2).

Fig. 5 depicts the comparison results, showing that 2-Gating is often inferior to *LassoNet* and *HSIC*. 3- and 4-Gating, however, dominate all other methods for almost all possible input sparsity configurations and, hence, are the favorable options among these methods for feature selection.

5.2.2 Filter sparsity in convolutional neural networks

In the next experiment, we investigate filter sparsification — one of the most prominent applications of structured sparsity in neural networks. As comparison methods, we select three commonly used methods in the filter sparsity literature [19]: Global magnitude pruning, $L_{2,1}$ -penalization with naïve optimization followed by magnitude pruning (MP) [60], and network slimming [38]. A more detailed description can be found in Appendix C.3. We run experiments on CIFAR-10, CIFAR-100, and SVHN, using a VGG-16 [49] and ResNet-18 model [18]. *D*-Gating is implemented by adding gating parameters on the filter level, which, given the size of these models, has a negligible parameter overhead (see Table 5). As filter sparsity can be used to construct a smaller model, potentially deployable on edge devices or similar, we also measure the theoretical speed-up of the sparsified model using floating-point operations (FLOPs). Similar to previous results, Fig. 6 unveils a superior performance of $D > 2$ -Gating compared to gating with $D = 2$. However, all gating approaches outperform the filter sparsity baselines despite our approach not requiring post-hoc pruning as the main sparsification mechanism. This is the case both in terms of the accuracy-sparsity tradeoff provided by our method as well as the theoretical speed-up implied (cf. Table 1 and Fig. 11).




Figure 6: Structured accuracy-sparsity tradeoffs of *D*-Gated neural networks and comparison methods.

Table 1: Largest theoretical speed-up ($\frac{\text{base}}{\text{sparse}}$ FLOPs) within 5% / 10% of the max. test accuracy.

Method	VGG-16 (CIFAR-10)	VGG-16 (SVHN)	ResNet-18 (CIFAR-100)
D-Gating ($D = 2$)	16.13 / 52.70	175.49 / 175.49	5.13 / 14.91
D-Gating ($D = 3$)	11.17 / 25.51	266.66 / 390.03	8.27 / 15.14
D-Gating ($D = 4$)	18.88 / 48.97	262.05 / 541.65	12.76 / 41.04
$L_{2,1}$ + Mag. pruning	6.97 / 6.97	18.88 / 18.88	3.35 / 14.01
Network Slimming	2.45 / 4.91	7.53 / 7.53	3.56 / 3.56

5.2.3 Structured sparsity in language modeling

Our next application considers the effect of *D*-Gating in an attention-based language model [57]. For this, we use NanoGPT and apply *D*-Gating to the attention heads of all attention layers. A natural comparison is again the direct optimization of the $L_{2,1}$ -penalty, i.e., without first transforming the objective into a differentiable one through *D*-Gating. To highlight the shortcomings of this naïve approach, we also follow the direct optimization with an explicit pruning step. We train NanoGPT on TinyShakespeare (details in Appendix C.4) and evaluate the model’s validation accuracy as well as validation perplexity for different regularization strengths and hence levels of attention head sparsity.

Fig. 7 confirms our hypothesis that direct optimization does not result in structured sparsity. Notably, the min. and max. norms of the attention heads converge for large λ under direct $L_{2,1}$ penalization. In contrast, *D*-Gating shows the desired effect of increased sparsity for higher regularization and provides a smooth tradeoff between accuracy and sparsity. Even when combining direct optimization

of the $L_{2,1}$ regularized objective with additional post-hoc pruning, and taking, at each pruning ratio, the best performance over a grid of λ values, we see that D -Gating achieves much higher head sparsity values before performance degrades significantly.




Figure 7: Structured sparsity-performance tradeoffs for NanoGPT trained on TinyShakespeare with D -Gated Value matrices in their attention heads and direct $L_{2,1}$ regularization. **Left:** Val. accuracy vs. head sparsity. **Middle:** Val. perplexity vs. λ . **Right two:** Effect of regularization λ on head sparsity and evolution of 2-norms for $D \in \{2, 3, 4\}$. Dashed horizontal lines represent the vanilla performance.

5.2.4 Tree sparsity in neural trees

As a final application, we investigate the sparsification of neural trees. More specifically, we propose a novel modification of Neural Oblivious Decision Ensembles (NODE) [45], a neural network-based decision tree ensemble model. While there are multiple options to apply our approach within this architecture, we demonstrate the efficacy of D -Gating by inducing sparsity on the tree-level, i.e., using the different trees as groups. We test our approach on the Wine data set [6] using a range of λ values for a single tree-layer as suggested by [45]. The tree layer consists of 500 trees, and the weights corresponding to each tree are gated with $D = 2$ to induce differentiable $L_{2,1}$ group sparsity. This is already lower than the default hyperparameter for the tree count of 2048 reported in [45], but raises the question whether 500 trees are in fact necessary to obtain reported performances. Fig. 8 shows the test root mean squared error (RMSE) and the number of active trees as functions of λ . While the full non-sparse model reaches baseline RMSE values, we observe that it is possible to achieve a very similar performance by using less than 5 trees, e.g., for $\lambda = 3 \times 10^{-3}$. This suggests that, at least for the Wine data set, a much simpler and notably less expensive configuration is sufficient.




Figure 8: Regularization paths for NODE on the Wine dataset with 2-Gating on tree level. Means and standard deviations over four data splits are shown.

6 Discussion

In this paper, we introduce D -Gating, a differentiable structured sparsity method compatible with SGD and applicable to arbitrary differentiable architectures, addressing limitations of non-differentiable penalties in deep learning. We thereby positively answer our initial research question on whether it is possible to design a modular structured sparsity routine, integrable into any architecture, amenable to SGD, with theoretical sparsity guarantees and little practical overhead.

Limitations and future work Due to the flexibility of D -Gating, our approach can provide structured sparsity penalties for arbitrary grouping structures. We systematically demonstrate this flexibility through a diverse set of applications in Section 5.2 and Appendix E. While our theoretical results guarantee equivalence to the original sparse but non-smooth optimization problem, future work could further explore the benefits of this formulation or assess its performance when combined with sophisticated pruning and retraining schedules. Secondly, although the gradient flow limit admits clear analysis, it remains an open question how discrete-time SGD deviates in practice. Finally, integrating D -Gating into the complex training pipelines of modern large-scale foundation models, where sparse training from scratch is often impractical, presents another promising direction.

References

- [1] Davide Anguita, Alessandro Ghio, Luca Oneto, Xavier Parra, Jorge Luis Reyes-Ortiz, et al. A public domain dataset for human activity recognition using smartphones. In *Esann*, volume 3, pages 3–4, 2013.
- [2] Davis Blalock, Jose Javier Gonzalez Ortiz, Jonathan Frankle, and John Gutttag. What is the state of neural network pruning? *Proceedings of machine learning and systems*, 2:129–146, 2020.
- [3] Kevin Bui, Fredrick Park, Shuai Zhang, Yingyong Qi, and Jack Xin. Improving network slimming with nonconvex regularization. *IEEE Access*, 9:115292–115314, 2021.
- [4] Kevin Bui, Fredrick Park, Shuai Zhang, Yingyong Qi, and Jack Xin. Structured sparsity of convolutional neural networks via nonconvex sparse group regularization. *Frontiers in applied mathematics and statistics*, 6:529564, 2021.
- [5] Hongrong Cheng, Miao Zhang, and Javen Qinfeng Shi. A survey on deep neural network pruning: Taxonomy, comparison, analysis, and recommendations. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024.
- [6] P. Cortez, A. Cerdeira, F. Almeida, T. Matos, and J. Reis. Wine Quality. UCI Machine Learning Repository, 2009. DOI: <https://doi.org/10.24432/C56S3T>.
- [7] Tristan Deleu and Yoshua Bengio. Structured sparsity inducing adaptive optimizers for deep learning. *arXiv preprint arXiv:2102.03869*, 2021.
- [8] Mark Fanty and Ronald Cole. Spoken letter recognition. *Advances in neural information processing systems*, 3, 1990.
- [9] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Finding sparse, trainable neural networks. In *International Conference on Learning Representations*, 2019.
- [10] Jonathan Frankle, Gintare Karolina Dziugaite, Daniel Roy, and Michael Carbin. Pruning neural networks at initialization: Why are we missing the mark? In *International Conference on Learning Representations*, 2020.
- [11] Elias Frantar and Dan Alistarh. Sparsegpt: Massive language models can be accurately pruned in one-shot. In *International Conference on Machine Learning*, pages 10323–10337. PMLR, 2023.
- [12] Konrad Gadzicki, Razieh Khamsehashari, and Christoph Zetzsche. Early vs late fusion in multimodal convolutional neural networks. In *2020 IEEE 23rd international conference on information fusion (FUSION)*, pages 1–6. IEEE, 2020.
- [13] Trevor Gale, Erich Elsen, and Sara Hooker. The state of sparsity in deep neural networks. *arXiv preprint arXiv:1902.09574*, 2019.
- [14] Yves Grandvalet. Least absolute shrinkage is equivalent to quadratic penalization. In *ICANN 98: Proceedings of the 8th International Conference on Artificial Neural Networks, Skövde, Sweden, 2–4 September 1998*, pages 201–206. Springer, 1998.
- [15] Suriya Gunasekar, Jason D Lee, Daniel Soudry, and Nati Srebro. Implicit bias of gradient descent on linear convolutional networks. *Advances in Neural Information Processing Systems*, 31, 2018.
- [16] Isabelle Guyon. Design of experiments of the nips 2003 variable selection benchmark. In *NIPS 2003 workshop on feature extraction and feature selection*, volume 253, page 40, 2003.
- [17] Song Han, Jeff Pool, John Tran, and William Dally. Learning both weights and connections for efficient neural network. *Advances in neural information processing systems*, 28, 2015.
- [18] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 770–778, 2016.

- [19] Torsten Hoeﬂer, Dan Alistarh, Tal Ben-Nun, Nikoli Dryden, and Alexandra Peste. Sparsity in deep learning: Pruning and growth for eﬃcient inference and training in neural networks. *Journal of Machine Learning Research*, 22(241):1–124, 2021.
- [20] Peter D Hoff. Lasso, fractional norm and structured sparse estimation using a hadamard product parametrization. *Computational Statistics & Data Analysis*, 115:186–198, 2017.
- [21] Yaohua Hu, Chong Li, Kaiwen Meng, Jing Qin, and Xiaoqi Yang. Group sparse optimization via $\ell_{p,q}$ regularization. *The Journal of Machine Learning Research*, 18(1):960–1011, 2017.
- [22] Zehao Huang and Naiyan Wang. Data-driven sparse structure selection for deep neural networks. In *Proceedings of the European conference on computer vision (ECCV)*, pages 304–320, 2018.
- [23] Tom Jacobs and Rebekka Burkholz. Mask in the mirror: Implicit sparsification. In *International Conference on Machine Learning*. PMLR, 2025.
- [24] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [25] Seijin Kobayashi, Yassir Akram, and Johannes Von Oswald. Weight decay induces low-rank attention layers. *Advances in Neural Information Processing Systems*, 37:4481–4510, 2024.
- [26] Chris Kolb, Christian L Müller, Bernd Bischl, and David Rügamer. Smoothing the edges: a general framework for smooth optimization in sparse regularization using hadamard over-parametrization. *arXiv preprint arXiv:2307.03571*, 2023.
- [27] Chris Kolb, Bernd Bischl, and David Rügamer. Differentiable attention sparsity via structured Soft -gating. In *ICLR Workshop on Sparsity in LLMs (SLLM): Deep Dive into Mixture of Experts, Quantization, Hardware, and Inference*, 2025.
- [28] Chris Kolb, Tobias Weber, Bernd Bischl, and David Rügamer. Deep weight factorization: Sparse learning through the lens of artificial symmetries. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [29] Alex Krizhevsky, Geoffrey Hinton, et al. Learning multiple layers of features from tiny images. 2009.
- [30] Daniel Kunin, Javier Sagastuy-Brena, Surya Ganguli, Daniel LK Yamins, and Hidenori Tanaka. Neural mechanics: Symmetry and broken conservation laws in deep learning dynamics. *arXiv preprint arXiv:2012.04728*, 2020.
- [31] Aditya Kusupati, Vivek Ramanujan, Raghav Somani, Mitchell Wortsman, Prateek Jain, Sham Kakade, and Ali Farhadi. Soft threshold weight reparameterization for learnable sparsity. In *International Conference on Machine Learning*, pages 5544–5555. PMLR, 2020.
- [32] Yann LeCun, John Denker, and Sara Solla. Optimal brain damage. *Advances in neural information processing systems*, 2, 1989.
- [33] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- [34] N Lee, T Ajanthan, and P Torr. Snip: single-shot network pruning based on connection sensitivity. In *International Conference on Learning Representations*. Open Review, 2019.
- [35] Ismael Lemhadri, Feng Ruan, and Rob Tibshirani. Lassonet: Neural networks with feature sparsity. In *International Conference on Artificial Intelligence and Statistics*, pages 10–18. PMLR, 2021.
- [36] Hao Li, Asim Kadav, Igor Durdanovic, Hanan Samet, and Hans Peter Graf. Pruning filters for efficient convnets. In *International Conference on Learning Representations*, 2017.
- [37] Jiangyuan Li, Thanh V Nguyen, Chinmay Hegde, and Raymond KW Wong. Implicit regularization for group sparsity. In *The Eleventh International Conference on Learning Representations*, 2023.

- [38] Zhuang Liu, Jianguo Li, Zhiqiang Shen, Gao Huang, Shoumeng Yan, and Changshui Zhang. Learning efficient convolutional networks through network slimming. In *Proceedings of the IEEE international conference on computer vision*, pages 2736–2744, 2017.
- [39] Miao Lu, Xiaolong Luo, Tianlong Chen, Wuyang Chen, Dong Liu, and Zhangyang Wang. Learning pruning-friendly networks via frank-wolfe: One-shot, any-sparsity, and no retraining. In *International Conference on Learning Representations*, 2022.
- [40] Xinyin Ma, Gongfan Fang, and Xinchao Wang. Llm-pruner: On the structural pruning of large language models. *Advances in neural information processing systems*, 36:21702–21720, 2023.
- [41] Morteza Mardani, Qingyun Sun, David Donoho, Vardan Papyan, Hatef Monajemi, Shreyas Vasanaawala, and John Pauly. Neural proximal gradient descent for compressive imaging. *Advances in Neural Information Processing Systems*, 31, 2018.
- [42] Sameer A Nene, Shree K Nayar, Hiroshi Murase, et al. Columbia object image library (coil-20). 1996.
- [43] Yuval Netzer, Tao Wang, Adam Coates, Alessandro Bissacco, Baolin Wu, Andrew Y Ng, et al. Reading digits in natural images with unsupervised feature learning. In *NIPS workshop on deep learning and unsupervised feature learning*, volume 2011, page 4. Granada, 2011.
- [44] Alexandra Peste, Eugenia Iofinova, Adrian Vladu, and Dan Alistarh. Ac/dc: Alternating compressed/decompressed training of deep neural networks. *Advances in neural information processing systems*, 34:8557–8570, 2021.
- [45] Sergei Popov, Stanislav Morozov, and Artem Babenko. Neural oblivious decision ensembles for deep learning on tabular data. In *International Conference on Learning Representations*, 2020.
- [46] Pedro Savarese, Hugo Silva, and Michael Maire. Winning the lottery with continuous sparsification. *Advances in neural information processing systems*, 33:11380–11390, 2020.
- [47] Simone Scardapane, Danilo Comminiello, Amir Hussain, and Aurelio Uncini. Group sparse regularization for deep neural networks. *Neurocomputing*, 241:81–89, 2017.
- [48] Noah Simon, Jerome Friedman, Trevor Hastie, and Robert Tibshirani. A sparse-group lasso. *Journal of computational and graphical statistics*, 22(2):231–245, 2013.
- [49] K Simonyan and A Zisserman. Very deep convolutional networks for large-scale image recognition. In *3rd International Conference on Learning Representations (ICLR 2015)*. Computational and Biological Learning Society, 2015.
- [50] Nathan Srebro. *Learning with Matrix Factorizations*. PhD thesis, Massachusetts Institute of Technology, 2004.
- [51] Mingjie Sun, Zhuang Liu, Anna Bair, and J Zico Kolter. A simple and effective pruning approach for large language models. In *The Twelfth International Conference on Learning Representations*, 2023.
- [52] Hidenori Tanaka, Daniel Kunin, Daniel L Yamins, and Surya Ganguli. Pruning neural networks without any data by iteratively conserving synaptic flow. *Advances in neural information processing systems*, 33:6377–6389, 2020.
- [53] Yingjie Tian and Yuqi Zhang. A comprehensive survey on regularization strategies in machine learning. *Information Fusion*, 80:146–166, 2022.
- [54] Robert Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B (Methodological)*, 58(1):267–288, 1996.
- [55] Ryan Tibshirani. Equivalences between sparse models and neural networks. *Working Notes*, 2021.
- [56] Tomas Vaskevicius, Varun Kanade, and Patrick Rebeschini. Implicit regularization for optimal sparse recovery. *Advances in Neural Information Processing Systems*, 32, 2019.

- [57] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.
- [58] Elena Voita, David Talbot, Fedor Moiseev, Rico Sennrich, and Ivan Titov. Analyzing multi-head self-attention: Specialized heads do the heavy lifting, the rest can be pruned. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 5797–5808, 2019.
- [59] Chaoqi Wang, Guodong Zhang, and Roger Grosse. Picking winning tickets before training by preserving gradient flow. In *International Conference on Learning Representations*, 2020.
- [60] Wei Wen, Chunpeng Wu, Yandan Wang, Yiran Chen, and Hai Li. Learning structured sparsity in deep neural networks. In *Advances in Neural Information Processing Systems*, volume 29, 2016.
- [61] Blake Woodworth, Suriya Gunasekar, Jason D Lee, Edward Moroshko, Pedro Savarese, Itay Golan, Daniel Soudry, and Nathan Srebro. Kernel and rich regimes in overparametrized models. In *Conference on Learning Theory*, pages 3635–3673. PMLR, 2020.
- [62] Han Xiao, Kashif Rasul, and Roland Vollgraf. Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*, 2017.
- [63] Makoto Yamada, Wittawat Jitkrittum, Leonid Sigal, Eric P Xing, and Masashi Sugiyama. High-dimensional feature selection by feature-wise kernelized lasso. *Neural computation*, 26(1): 185–207, 2014.
- [64] Ming Yuan and Yi Lin. Model selection and estimation in regression with grouped variables. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 68(1):49–67, 2006.
- [65] Sergey Zagoruyko. 92.45% on cifar-10 in torch. <https://torch.ch/blog/2015/07/30/cifar.html>, 2015. Torch Blog.
- [66] Qiaozhe Zhang, Ruijie Zhang, Jun Sun, and Yingzhuang Liu. How sparse can we prune a deep network: A fundamental limit perspective. *Advances in Neural Information Processing Systems*, 37:91337–91372, 2024.
- [67] Peng Zhao, Yun Yang, and Qiao-Chu He. High-dimensional linear regression via implicit regularization. *Biometrika*, 2022.
- [68] Haizhong Zheng, Xiaoyan Bai, Xueshen Liu, Zhuoqing Morley Mao, Beidi Chen, Fan Lai, and Atul Prakash. Learn to be efficient: Build structured sparsity in large language models. *Advances in Neural Information Processing Systems*, 37:101969–101991, 2024.
- [69] Liu Ziyin. Symmetry induces structure and constraint of learning. In *Forty-first International Conference on Machine Learning*, 2023.
- [70] Liu Ziyin and Zihao Wang. spred: Solving l1 penalty with sgd. In *International Conference on Machine Learning*, pages 43407–43422. PMLR, 2023.

Appendices

A	Algorithm	15
B	Missing proofs	15
B.1	Proof of Lemma 1	15
B.2	Proof of Corollary 1	16
B.3	Proof of Theorem 1	17
B.4	Proof of Lemma 2	18
B.5	Proof of Lemma 3	19
B.6	Proof of Lemma 4	19
C	Experimental details	21
C.1	Details on group lasso simulation	21
C.2	Details on feature selection experiments	22
C.3	Details on filter sparsity experiments	22
C.4	Details on language modeling experiments	23
C.5	Details on Neural Oblivious Decision Ensembles	24
D	Computational Environment	24
E	Further experiments and ablation studies	24
E.1	Additional plots for filter sparsity in image classification	24
E.2	Overheads	24
E.3	Depth and instability ablation studies	25
E.4	Imbalance decay and loss convergence for SGD and Adam	26
E.5	Subnetwork selection	26

A Algorithm

In Algorithm 1, we provide the algorithm for sparse training using the proposed D -Gating method.

Algorithm 1 D -Gating Network Training

- 1: **Input:**
 - 2: Dataset $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$, architecture $\mathcal{A}$ with weights $\mathbf{w} \in \mathbb{R}^p$ partitioned into J groups $\{\mathcal{G}_j\}_{j=1}^J$
 - 3: Gating depth $D \geq 2$
 - 4: Training hyperparameters $\{T, |B|, \{\eta^{(t)}\}_{t=0}^T, \lambda\}$
 - 5: Threshold $\varepsilon_{\text{tiny}}$
 - 6: **D -Gating parametrization:**
 - 7: Replace weights $\mathbf{w}$ in $\mathcal{A}$ by primary weights $\boldsymbol{\omega} = (\boldsymbol{\omega}_j)_{j=1}^J$ and gating factors $\boldsymbol{\Gamma} = (\gamma_{j,d})_{j \in [J], d \in [D-1]}$
 - 8: Effective weights $\mathbf{w} = \boldsymbol{\omega} \blacktriangleright \boldsymbol{\gamma}^\odot$, so $\mathbf{w}_j = \boldsymbol{\omega}_j \prod_{d=1}^{D-1} \gamma_{j,d}$
 - 9: **Initialize** $\boldsymbol{\omega}^{(0)} \leftarrow \text{Standard-Init}(\mathbf{w}|\mathcal{A})$ and $\boldsymbol{\Gamma}^{(0)} \leftarrow \mathbf{1}_J \mathbf{1}_{D-1}^\top \in \mathbb{R}^{J \times (D-1)}$
 - 10: **for** $t = 0$ to $T - 1$ **do**
 - 11: Sample mini-batch $B^{(t)} \subseteq \mathcal{D}$
 - 12: Compute $\mathbf{w}^{(t)} = \boldsymbol{\omega}^{(t)} \blacktriangleright (\boldsymbol{\gamma}^\odot)^{(t)}$ and mini-batch loss
$$\mathcal{L}(\boldsymbol{\omega}^{(t)}, \boldsymbol{\Gamma}^{(t)}) = \sum_{(\mathbf{x}, y) \in B^{(t)}} \ell(y, f(\mathbf{x}; \mathbf{w}^{(t)})) + \frac{\lambda}{D} (\|\boldsymbol{\omega}^{(t)}\|_2^2 + \|\boldsymbol{\Gamma}^{(t)}\|_F^2).$$
 - 13: Compute gradients $\nabla_{\boldsymbol{\omega}} \mathcal{L}(\boldsymbol{\omega}, \boldsymbol{\Gamma}), \nabla_{\boldsymbol{\Gamma}} \mathcal{L}(\boldsymbol{\omega}, \boldsymbol{\Gamma})$
 - 14: Update
$$\boldsymbol{\omega}^{(t+1)} \leftarrow \boldsymbol{\omega}^{(t)} - \frac{\eta^{(t)}}{|B|} \nabla_{\boldsymbol{\omega}} \mathcal{L}(\boldsymbol{\omega}, \boldsymbol{\Gamma}), \quad \boldsymbol{\Gamma}^{(t+1)} \leftarrow \boldsymbol{\Gamma}^{(t)} - \frac{\eta^{(t)}}{|B|} \nabla_{\boldsymbol{\Gamma}} \mathcal{L}(\boldsymbol{\omega}, \boldsymbol{\Gamma}).$$
 - 15: $\eta^{(t+1)} \leftarrow \text{LRSchedule}(t + 1)$
 - 16: **end for**
 - 17: **Collapse gates and reduce architecture:**
 - 18: $\hat{\mathbf{w}} = \boldsymbol{\omega}^{(T)} \blacktriangleright (\boldsymbol{\gamma}^\odot)^{(T)}$
 - 19: Zero out removed weight groups: $\hat{\mathbf{w}}_j \leftarrow \mathbf{0}$ if $\|\hat{\mathbf{w}}_j\|_2 < \varepsilon_{\text{tiny}}$
 - 20: Load sparse weights $\hat{\mathbf{w}}$ into compact architecture $\tilde{\mathcal{A}}$
 - 21: **Output:** Sparse network parameters $\hat{\mathbf{w}}$ and reduced architecture $\tilde{\mathcal{A}}$
-

B Missing proofs

B.1 Proof of Lemma 1

Proof. We argue by contradiction. Suppose that for some group j the gating factors are not balanced; that is, there exist indices $d, d' \in [D - 1]$ such that, without loss of generality,

$$\gamma_{j,d}^2 < \gamma_{j,d'}^2. \quad (17)$$

An analogous argument applies if one considers an imbalance between $\|\boldsymbol{\omega}_j\|_2^2$ and one of the scalar factors. We now apply a first-order perturbation argument. Consider an infinitesimal perturbation parameter ε and define the perturbed factors as

$$\tilde{\gamma}_{j,d} = \gamma_{j,d}(1 + \varepsilon) \quad \text{and} \quad \tilde{\gamma}_{j,d'} = \gamma_{j,d'}(1 - \varepsilon). \quad (18)$$

Noting that

$$(1 + \varepsilon)^2 \approx 1 + 2\varepsilon, \quad (1 - \varepsilon)^2 \approx 1 - 2\varepsilon, \quad (19)$$

the product of the perturbed factors satisfies

$$\begin{aligned} \tilde{\gamma}_{j,d} \tilde{\gamma}_{j,d'} &= \gamma_{j,d}(1 + \varepsilon) \gamma_{j,d'}(1 - \varepsilon) \\ &= \gamma_{j,d} \gamma_{j,d'} (1 + \varepsilon)(1 - \varepsilon) \\ &= \gamma_{j,d} \gamma_{j,d'} (1 - \varepsilon^2) \approx \gamma_{j,d} \gamma_{j,d'}, \end{aligned} \quad (20)$$

so that the effective parameter $\mathbf{w}_j = \boldsymbol{\omega}_j \prod_{d=1}^{D-1} \gamma_{j,d}$ remains unchanged to first order. On the other hand, the original L_2 penalty for the two factors is given by

$$\mathcal{R}(\gamma_{j,d}, \gamma_{j,d'}) = \gamma_{j,d}^2 + \gamma_{j,d'}^2. \quad (21)$$

After perturbation, we have

$$(\tilde{\gamma}_{j,d})^2 \approx \gamma_{j,d}^2(1 + 2\varepsilon), \quad (\tilde{\gamma}_{j,d'})^2 \approx \gamma_{j,d'}^2(1 - 2\varepsilon), \quad (22)$$

so that the new penalty becomes

$$\begin{aligned} \mathcal{R}(\tilde{\gamma}_{j,d}, \tilde{\gamma}_{j,d'}) &\approx \gamma_{j,d}^2(1 + 2\varepsilon) + \gamma_{j,d'}^2(1 - 2\varepsilon) \\ &= (\gamma_{j,d}^2 + \gamma_{j,d'}^2) + 2\varepsilon(\gamma_{j,d}^2 - \gamma_{j,d'}^2). \end{aligned} \quad (23)$$

Since $\gamma_{j,d}^2 < \gamma_{j,d'}^2$ by Eq. (17), it follows that

$$2\varepsilon(\gamma_{j,d}^2 - \gamma_{j,d'}^2) < 0. \quad (24)$$

This strict decrease in the L_2 penalty contradicts the stationarity of $(\boldsymbol{\omega}, \boldsymbol{\Gamma})$ with respect to $\mathcal{L}(\boldsymbol{\omega}, \boldsymbol{\Gamma})$. Therefore, the gating factors must be balanced at any stationary point, which, together with the invariance of the effective parameter $\mathbf{w}$, implies Eq. (7). $\square$

B.2 Proof of Corollary 1

Proof. We begin by group-wise separating the smooth surrogate penalty in (5):

$$\frac{1}{D}(\|\boldsymbol{\omega}\|_2^2 + \|\boldsymbol{\Gamma}\|_F^2) = \sum_{j=1}^J \frac{1}{D} \left(\|\boldsymbol{\omega}_j\|_2^2 + \sum_{d=1}^{D-1} \gamma_{j,d}^2 \right). \quad (25)$$

Next, for each group $j \in [J]$, we apply the inequality of arithmetic and geometric means (AM-GM) to the D non-negative terms $\|\boldsymbol{\omega}_j\|_2^2, \gamma_{j,1}^2, \dots, \gamma_{j,D-1}^2$:

$$\frac{1}{D} \left(\|\boldsymbol{\omega}_j\|_2^2 + \sum_{d=1}^{D-1} \gamma_{j,d}^2 \right) \geq (\|\boldsymbol{\omega}_j\|_2^2 \cdot \prod_{d=1}^{D-1} \gamma_{j,d}^2)^{1/D}. \quad (26)$$

Since $\mathbf{w}_j = \boldsymbol{\omega}_j \prod_{d=1}^{D-1} \gamma_{j,d}$, the right-hand side becomes

$$(\|\boldsymbol{\omega}_j\|_2^2 \cdot \prod_{d=1}^{D-1} \gamma_{j,d}^2)^{1/D} = |\prod_{d=1}^{D-1} \gamma_{j,d} \cdot \|\boldsymbol{\omega}_j\|_2|^{2/D} = \|\mathbf{w}_j\|_2^{2/D} = \|\mathbf{w}_j\|_2^{2/D}. \quad (27)$$

Summing over all J groups yields

$$\frac{1}{D}(\|\boldsymbol{\omega}\|_2^2 + \|\boldsymbol{\Gamma}\|_F^2) \geq \sum_{j=1}^J \|\mathbf{w}_j\|_2^{2/D} = \|\mathbf{w}\|_{2,2/D}^{2/D}, \quad (28)$$

with equality holding if and only if $\|\boldsymbol{\omega}_j\|_2^2 = \gamma_{j,1}^2 = \dots = \gamma_{j,D-1}^2$ for each $j \in [J]$, i.e. exactly the balancedness condition (7). Under that condition, each group-wise penalty attains its minimal value $\|\mathbf{w}_j\|_2^{2/D}$ subject to the constraint that the effective parameter $\mathbf{w}_j$ remains unchanged. Finally, substituting back into $\mathcal{L}(\boldsymbol{\omega}, \boldsymbol{\Gamma}) = \mathcal{L}_0(\boldsymbol{\omega} \blacktriangleright \boldsymbol{\gamma}^\odot) + \frac{\lambda}{D}(\|\boldsymbol{\omega}\|_2^2 + \|\boldsymbol{\Gamma}\|_F^2)$ gives

$$\mathcal{L}(\boldsymbol{\omega}, \boldsymbol{\Gamma}) = \mathcal{L}_0(\mathbf{w}) + \lambda \sum_{j=1}^J \|\mathbf{w}_j\|_2^{2/D} = \mathcal{L}_0(\mathbf{w}) + \lambda \|\mathbf{w}\|_{2,2/D}^{2/D} = \mathcal{L}_{\mathbf{w}}(\mathbf{w}), \quad (29)$$

completing the proof. $\square$

B.3 Proof of Theorem 1

Proof. The two objectives are

$$\mathcal{L}(\omega, \Gamma) = \mathcal{L}_0(\omega \blacktriangleright \gamma^\odot) + \frac{\lambda}{D} (\|\omega\|_2^2 + \|\Gamma\|_F^2), \quad (30)$$

$$\mathcal{L}_{\mathbf{w}}(\mathbf{w}) = \mathcal{L}_0(\mathbf{w}) + \lambda \sum_{j=1}^J \|\mathbf{w}_j\|_2^{2/D}. \quad (31)$$

For the first direction, suppose that $\hat{\mathbf{w}}$ is a local minimizer of $\mathcal{L}_{\mathbf{w}}(\mathbf{w})$. Then there is $\varepsilon > 0$ such that

$$\mathcal{L}_{\mathbf{w}}(\hat{\mathbf{w}}) \leq \mathcal{L}_{\mathbf{w}}(\mathbf{w}) \quad \forall \mathbf{w} \in \mathcal{B}(\hat{\mathbf{w}}, \varepsilon), \quad (32)$$

where $\mathcal{B}(\hat{\mathbf{w}}, \varepsilon)$ denotes an ε -ball around $\hat{\mathbf{w}} \in \mathbb{R}^p$. By the multiplicative and surjective nature of the D -Gating overparametrization $\blacktriangleright$, we can always choose a *balanced* representation $(\hat{\omega}, \hat{\Gamma})$ in the preimage of $\hat{\mathbf{w}}$, i.e.,

$$\hat{\omega} \blacktriangleright \hat{\gamma}^\odot = \hat{\mathbf{w}}, \quad \|\hat{\omega}_j\|_2^2 = \hat{\gamma}_{j,1}^2 = \dots = \|\hat{\mathbf{w}}_j\|_2^{2/D} \quad \forall j = 1, \dots, J. \quad (33)$$

Using Corollary 1, balancedness implies that the D -gated objective simplifies to

$$\mathcal{L}(\hat{\omega}, \hat{\Gamma}) = \mathcal{L}_{\mathbf{w}}(\hat{\mathbf{w}}). \quad (34)$$

Further, by continuity of $\blacktriangleright$, there is $\delta > 0$ such that

$$\|(\omega, \Gamma) - (\hat{\omega}, \hat{\Gamma})\|_2 < \delta \implies \|\omega \blacktriangleright \gamma^\odot - \hat{\mathbf{w}}\|_2 < \varepsilon. \quad (35)$$

Hence all $(\omega, \Gamma) \in \mathcal{B}((\hat{\omega}, \hat{\Gamma}), \delta)$ map to some effective weight $\mathbf{w} = \omega \blacktriangleright \gamma^\odot \in \mathcal{B}(\hat{\mathbf{w}}, \varepsilon)$. Moreover, the gated objective $\mathcal{L}(\omega, \Gamma)$ can be related to its non-differentiable counterpart $\mathcal{L}_{\mathbf{w}}(\mathbf{w})$ as $\mathcal{L}(\omega, \Gamma) = \mathcal{L}_{\mathbf{w}}(\omega \blacktriangleright \gamma^\odot) + \lambda \mathcal{M}(\omega, \Gamma)$, where $\mathcal{M}(\omega, \Gamma) := D^{-1}(\|\omega\|_2^2 + \|\Gamma\|_F^2) - \|\omega \blacktriangleright \gamma^\odot\|_2^{2/D}$ measures the (non-negative) misalignment of the D -Gating variables, achieving 0 if and only if (ω, Γ) is a balanced representation of the effective weight $\mathbf{w} = \omega \blacktriangleright \gamma^\odot$. Then, using Eq. (32) and Eq. (34), we obtain the following chain of inequalities

$$\begin{aligned} \forall (\omega, \Gamma) \in \mathcal{B}((\hat{\omega}, \hat{\Gamma}), \delta) : \mathcal{L}(\omega, \Gamma) &= \mathcal{L}_{\mathbf{w}}(\omega \blacktriangleright \gamma^\odot) + \underbrace{\lambda \mathcal{M}(\omega, \Gamma)}_{\geq 0} \geq \mathcal{L}_{\mathbf{w}}(\omega \blacktriangleright \gamma^\odot) \\ &\geq \mathcal{L}_{\mathbf{w}}(\hat{\mathbf{w}}) \\ &= \mathcal{L}(\hat{\omega}, \hat{\Gamma}), \end{aligned} \quad (36)$$

showing $(\hat{\omega}, \hat{\Gamma})$ is a local minimizer of $\mathcal{L}(\omega, \Gamma)$.

To show the reverse direction, assume $(\hat{\omega}, \hat{\Gamma})$ is a local minimizer of $\mathcal{L}(\omega, \Gamma)$. We can apply Lemma 1 to establish the balancedness of the gating parameters $(\hat{\omega}, \hat{\Gamma})$ and further define the effective weight as $\hat{\mathbf{w}} = \hat{\omega} \blacktriangleright \hat{\gamma}^\odot$. Since $(\hat{\omega}, \hat{\Gamma})$ is balanced, $\mathcal{M}(\hat{\omega}, \hat{\Gamma}) = 0$ and thus $\mathcal{L}(\hat{\omega}, \hat{\Gamma}) = \mathcal{L}_{\mathbf{w}}(\hat{\mathbf{w}})$. By local minimality of $(\hat{\omega}, \hat{\Gamma})$, there exists $\delta > 0$ such that

$$\mathcal{L}(\hat{\omega}, \hat{\Gamma}) \leq \mathcal{L}(\omega, \Gamma) \quad \forall (\omega, \Gamma) \in \mathcal{B}((\hat{\omega}, \hat{\Gamma}), \delta). \quad (37)$$

Define for each group $j \in [J]$ the continuous function $r_j(\mathbf{w}_j) = \|\mathbf{w}_j\|_2^{1/D}$ on $\mathbb{R}^{p_j}$. Given any $\mathbf{w}' \in \mathbb{R}^p$ in a neighborhood of $\hat{\mathbf{w}}$, we can construct corresponding balanced D -Gating parameters as follows:

$$\gamma'_{j,d} = \begin{cases} \text{sign}(\hat{\gamma}_{j,d}) r_j(\mathbf{w}'_j), & \hat{\gamma}_{j,d} \neq 0, \\ r_j(\mathbf{w}'_j), & \hat{\gamma}_{j,d} = 0. \end{cases} \quad \omega'_j = \begin{cases} \mathbf{w}'_j / \prod_{d=1}^{D-1} \gamma'_{j,d}, & r_j(\mathbf{w}'_j) \neq 0, \\ \mathbf{0}, & r_j(\mathbf{w}'_j) = 0. \end{cases} \quad (38)$$

Then $\omega'_j \prod_{d=1}^{D-1} \gamma'_{j,d} = \mathbf{w}'_j$, and by construction, $\|\omega'_j\|_2^2 = \gamma'^2_{j,d} = r_j(\mathbf{w}'_j)^2 = \|\mathbf{w}'_j\|_2^{2/D}$ for all $j \in [J]$ and $d \in [D-1]$. Let (ω', Γ') denote the collection of group-wise gating parameters, i.e., $(\omega', \Gamma') = (\{\omega'_j\}_{j=1}^J, \{\gamma'_{j,d}\}_{j=1}^J, \{d=1\}^{D-1})$. By Corollary 1, we obtain $\mathcal{L}(\omega', \Gamma') = \mathcal{L}_{\mathbf{w}}(\mathbf{w}')$. Continuity of r_j and of the map $\mathbf{w}'_j \mapsto \omega'_j$ guarantees the existence of $\varepsilon > 0$ such that

$$\|\mathbf{w}' - \hat{\mathbf{w}}\|_2 < \varepsilon \implies \|(\omega', \Gamma') - (\hat{\omega}, \hat{\Gamma})\|_2 < \delta. \quad (39)$$

Therefore, for every $\mathbf{w}' \in \mathcal{B}(\hat{\mathbf{w}}, \varepsilon)$, the constructed D -Gating parameters (ω', Γ') satisfy Eq. (37) and thus

$$\forall \mathbf{w}' \in \mathcal{B}(\hat{\mathbf{w}}, \varepsilon) : \mathcal{L}_{\mathbf{w}}(\mathbf{w}') = \mathcal{L}(\omega', \Gamma') \geq \mathcal{L}(\hat{\omega}, \hat{\Gamma}) = \mathcal{L}_{\mathbf{w}}(\hat{\mathbf{w}}).$$

It follows that $\hat{\mathbf{w}}$ is a local minimizer of $\mathcal{L}_{\mathbf{w}}(\mathbf{w})$, completing the proof. $\square$

B.4 Proof of Lemma 2

Proof. We differentiate the pair-wise imbalance defined in Eq. (13) with respect to time and treat the two cases of $d = 1$ and $d \neq 1$ separately. In the following we use $\gamma_j^\odot \in \mathbb{R}$ to abbreviate $\prod_{d=1}^{D-1} \gamma_{j,d}$.

Case 1: For $d = 1$, we have

$$\mathcal{I}_{j,1,d'}(t) = \|\boldsymbol{\omega}_j(t)\|_2^2 - \gamma_{j,d'}(t)^2.$$

Differentiating with respect to time yields

$$\frac{d}{dt} \mathcal{I}_{j,1,d'}(t) = 2 \boldsymbol{\omega}_j(t)^\top \dot{\boldsymbol{\omega}}_j(t) - 2 \gamma_{j,d'}(t) \dot{\gamma}_{j,d'}(t). \quad (40)$$

Substituting the dynamics from Eq. (12) (with the appropriate negative signs) gives

$$\begin{aligned} \frac{d}{dt} \mathcal{I}_{j,1,d'}(t) = & -2 \left[\gamma_j^\odot \boldsymbol{\omega}_j(t)^\top \nabla_{\mathbf{w}_j} \mathcal{L}_0 + \frac{2\lambda}{D} \|\boldsymbol{\omega}_j(t)\|_2^2 \right] \\ & + 2 \gamma_{j,d'}(t) \left[\left(\frac{\gamma_j^\odot}{\gamma_{j,d'}(t)} \right) \boldsymbol{\omega}_j(t)^\top \nabla_{\mathbf{w}_j} \mathcal{L}_0 + \frac{2\lambda}{D} \gamma_{j,d'}(t) \right]. \end{aligned} \quad (41)$$

Since

$$\gamma_j^\odot = \gamma_{j,d'}(t) \left(\frac{\gamma_j^\odot}{\gamma_{j,d'}(t)} \right),$$

the terms involving $\boldsymbol{\omega}_j(t)^\top \nabla_{\mathbf{w}_j} \mathcal{L}_0$ cancel, leaving

$$\frac{d}{dt} \mathcal{I}_{j,1,d'}(t) = -\frac{4\lambda}{D} \left(\|\boldsymbol{\omega}_j(t)\|_2^2 - \gamma_{j,d'}(t)^2 \right) = -\frac{4\lambda}{D} \mathcal{I}_{j,1,d'}(t). \quad (42)$$

Case 2: For $d \neq 1$, we have

$$\mathcal{I}_{j,d,d'}(t) = \gamma_{j,d}(t)^2 - \gamma_{j,d'}(t)^2.$$

Differentiating with respect to t yields

$$\frac{d}{dt} \mathcal{I}_{j,d,d'}(t) = 2 \gamma_{j,d}(t) \dot{\gamma}_{j,d}(t) - 2 \gamma_{j,d'}(t) \dot{\gamma}_{j,d'}(t). \quad (43)$$

Substituting the expressions from Eq. (12) for both $\dot{\gamma}_{j,d}(t)$ and $\dot{\gamma}_{j,d'}(t)$ and using the identity

$$\gamma_j^\odot = \gamma_{j,d}(t) \left(\frac{\gamma_j^\odot}{\gamma_{j,d}(t)} \right),$$

the cancellation of the terms involving $\boldsymbol{\omega}_j(t)^\top \nabla_{\mathbf{w}_j} \mathcal{L}_0$ follows analogously to the first case, yielding

$$\frac{d}{dt} \mathcal{I}_{j,d,d'}(t) = -\frac{4\lambda}{D} \left(\gamma_{j,d}(t)^2 - \gamma_{j,d'}(t)^2 \right) = -\frac{4\lambda}{D} \mathcal{I}_{j,d,d'}(t). \quad (44)$$

Thus, in both cases, we have

$$\frac{d}{dt} \mathcal{I}_{j,d,d'}(t) = -\frac{4\lambda}{D} \mathcal{I}_{j,d,d'}(t),$$

so that the solution to the differential equation is

$$\mathcal{I}_{j,d,d'}(t) = \mathcal{I}_{j,d,d'}(0) e^{-\frac{4\lambda}{D} t}.$$

□

B.5 Proof of Lemma 3

Proof. For a fixed group j , the squared gating parameters are denoted as

$$a_1(t) = \|\omega_j(t)\|_2^2, \quad a_{d+1}(t) = \gamma_{j,d}(t)^2, \quad d \in [D-1], \quad (45)$$

the group-level regularizer is $\mathcal{R}_j(t) := \frac{1}{D} \sum_{d=1}^D a_d(t)$, and the effective group weight is $\mathbf{w}_j(t) = \omega_j(t) \prod_{d=1}^{D-1} \gamma_{j,d}(t) = \omega_j(t) \cdot \gamma_j^\odot(t)$, so that by construction $\|\mathbf{w}_j(t)\|_2^{2/D} = \left(\prod_{d=1}^D a_d(t)\right)^{1/D}$. Note that by Lemma 2, for each group $j \in [J]$, the pair-wise imbalances (as in Eq. (13)) satisfy

$$\mathcal{I}_{j,d,d'}(t) = \mathcal{I}_{j,d,d'}(0)e^{-\frac{4\lambda}{D}t}. \quad (46)$$

The overall misalignment in group j , measuring the balancedness of the gating representation, is given by

$$\mathcal{M}_j(t) := \mathcal{R}_j(t) - \|\mathbf{w}_j(t)\|_2^{2/D} \geq 0. \quad (47)$$

Let $m_j(t) = \min_{1 \leq d \leq D} a_d(t)$ and $M_j(t) = \max_{1 \leq d \leq D} a_d(t)$. Then, by definition, the maximal pair-wise imbalance for group $j \in [J]$ is $\mathcal{I}_{\max,j}(t) := M_j(t) - m_j(t)$. Since $\mathcal{R}_j(t)$ is an arithmetic mean, it satisfies $m_j(t) \leq \mathcal{R}_j(t) \leq M_j(t)$ and hence is upper bounded by $M_j(t)$. Because the geometric mean is not smaller than its smallest component, it is lower bounded by $m_j(t)$, i.e., $\|\mathbf{w}_j(t)\|_2^{2/D} = \left(\prod_{d=1}^D a_d(t)\right)^{1/D} \geq m_j(t)$, so that we have

$$\mathcal{M}_j(t) = \mathcal{R}_j(t) - \|\mathbf{w}_j(t)\|_2^{2/D} \leq M_j(t) - m_j(t) = \mathcal{I}_{\max,j}(t). \quad (48)$$

By Lemma 2, for any pair (d, d') in group j we have

$$|a_d(t) - a_{d'}(t)| \leq |a_d(0) - a_{d'}(0)|e^{-\frac{4\lambda}{D}t}, \quad (49)$$

so that $\mathcal{I}_{\max,j}(t) \leq \mathcal{I}_{\max,j}(0)e^{-\frac{4\lambda}{D}t}$. Hence, $\mathcal{M}_j(t) \leq \mathcal{I}_{\max,j}(0)e^{-\frac{4\lambda}{D}t}$ for every group j . Defining $\mathcal{I}_{\max}(0) = \max_{j \in \{1, \dots, J\}} \mathcal{I}_{\max,j}(0)$, we obtain $\mathcal{M}_j(t) \leq \mathcal{I}_{\max}(0)e^{-\frac{4\lambda}{D}t}$ for all j . Summing over j , the overall misalignment is

$$\mathcal{M}(t) = \sum_{j=1}^J \mathcal{M}_j(t) \leq J \mathcal{I}_{\max}(0)e^{-\frac{4\lambda}{D}t}. \quad (50)$$

By adding and subtracting $\mathcal{R}_{\mathbf{w}}(\mathbf{w}) := \|\mathbf{w}\|_{2,2/D}^{2/D}$ to the D -Gated objective $\mathcal{L}(\omega, \Gamma)$, we can express it as $\mathcal{L}_{\mathbf{w}}(\mathbf{w})$ plus the misalignment $\mathcal{M}(\omega, \Gamma) \geq 0$, i.e.,

$$\mathcal{L}(\omega, \Gamma) = \mathcal{L}_0(\omega \blacktriangleright \gamma^\odot) + \lambda \mathcal{R}(\omega, \Gamma) + \lambda (\mathcal{R}_{\mathbf{w}}(\mathbf{w}) - \mathcal{R}_{\mathbf{w}}(\mathbf{w})) = \mathcal{L}_{\mathbf{w}}(\omega \blacktriangleright \gamma^\odot) + \lambda \mathcal{M}(\omega, \Gamma). \quad (51)$$

As the difference between both losses is simply $\mathcal{L}(\omega(t), \Gamma(t)) - \mathcal{L}_{\mathbf{w}}(\mathbf{w}(t)) = \lambda \mathcal{M}(t)$, it follows that

$$\mathcal{L}(\omega(t), \Gamma(t)) - \mathcal{L}_{\mathbf{w}}(\mathbf{w}(t)) \leq \underbrace{\lambda J \mathcal{I}_{\max}(0)}_{:=C} e^{-\frac{4\lambda}{D}t} = C e^{-\frac{4\lambda}{D}t}, \quad (52)$$

where C is an initialization-dependent constant. We conclude that as $t \rightarrow \infty$, $\mathcal{L}(\omega(t), \Gamma(t))$ converges at least exponentially fast to $\mathcal{L}_{\mathbf{w}}(\mathbf{w}(t))$, with their difference vanishing with rate increasing in λ and decreasing in D . $\square$

B.6 Proof of Lemma 4

Proof. We first abbreviate the gradient of $\mathcal{L}_0$ w.r.t. $\mathbf{w}_j$ as $\mathbf{g}_j^{(t)} := \nabla_{\mathbf{w}_j} \mathcal{L}_0(\mathbf{w}^{(t)})$. For clarity of exposition, w.l.o.g, consider full-batch gradient descent updates, given by

$$\omega_j^{(t+1)} = \omega_j^{(t)} - \eta \left((\gamma_j^\odot)^{(t)} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \omega_j^{(t)} \right), \quad \gamma_{j,d}^{(t+1)} = \gamma_{j,d}^{(t)} - \eta \left(\prod_{d' \neq d} \gamma_{j,d'}^{(t)} \cdot \omega_j^{(t)\top} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \gamma_{j,d}^{(t)} \right), \quad (53)$$

and pair-wise imbalances $\mathcal{I}_{j,1,d'}^{(t)} = \|\omega_j^{(t)}\|_2^2 - (\gamma_{j,d'}^{(t)})^2$ and $\mathcal{I}_{j,d,d'}^{(t)} = (\gamma_{j,d}^{(t)})^2 - (\gamma_{j,d'}^{(t)})^2$ for $d, d' > 1$. We first inspect the vector-scalar case $\mathcal{I}_{j,1,d'}^{(t+1)}$ and start by expanding the squared norm of $\omega_j^{(t+1)}$:

$$\begin{aligned} \|\omega_j^{(t+1)}\|_2^2 &= \left\| \omega_j^{(t)} - \eta \left((\gamma_j^\odot)^{(t)} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \omega_j^{(t)} \right) \right\|_2^2 \\ &= \|\omega_j^{(t)}\|_2^2 - 2\eta (\gamma_j^\odot)^{(t)} \omega_j^{(t)\top} \mathbf{g}_j^{(t)} - \frac{4\lambda}{D} \eta \|\omega_j^{(t)}\|_2^2 + \eta^2 \left\| (\gamma_j^\odot)^{(t)} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \omega_j^{(t)} \right\|_2^2. \end{aligned} \quad (54)$$

Similarly, we expand the squared updated scalar $\gamma_{j,d'}^{(t+1)}$ and separate the terms by their order in η ,

$$\begin{aligned} (\gamma_{j,d'}^{(t+1)})^2 &= (\gamma_{j,d'}^{(t)})^2 - 2\eta \gamma_{j,d'}^{(t)} \cdot \prod_{d'' \neq d'} \gamma_{j,d''}^{(t)} \cdot \omega_j^{(t)\top} \mathbf{g}_j^{(t)} - \frac{4\lambda}{D} \eta (\gamma_{j,d'}^{(t)})^2 \\ &\quad + \eta^2 \left(\prod_{d'' \neq d'} \gamma_{j,d''}^{(t)} \cdot \omega_j^{(t)\top} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \gamma_{j,d'}^{(t)} \right)^2. \end{aligned} \quad (55)$$

Subtracting (55) from (54) to get $\mathcal{I}_{j,1,d'}^{(t+1)}$, the first-order terms $-2\eta (\gamma_j^\odot)^{(t)} \omega_j^{(t)\top} \mathbf{g}_j^{(t)}$ from both expansions cancel, since $(\gamma_j^\odot)^{(t)} = \gamma_{j,d'}^{(t)} \cdot \prod_{d'' \neq d'} \gamma_{j,d''}^{(t)}$, and we obtain after re-grouping of terms:

$$\mathcal{I}_{j,1,d'}^{(t+1)} = \left(1 - \frac{4\lambda}{D} \eta \right) \mathcal{I}_{j,1,d'}^{(t)} + \eta^2 \Delta_{j,1,d'}^{(t)}, \quad (56)$$

where

$$\begin{aligned} \Delta_{j,1,d'}^{(t)} &:= \left\| (\gamma_j^\odot)^{(t)} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \omega_j^{(t)} \right\|_2^2 - \left(\prod_{d'' \neq d'} \gamma_{j,d''}^{(t)} \cdot \omega_j^{(t)\top} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \gamma_{j,d'}^{(t)} \right)^2 \\ &= \|\nabla_{\omega_j} \mathcal{L}(\omega, \Gamma)\|_2^2 - \left(\frac{\partial \mathcal{L}(\omega, \Gamma)}{\partial \gamma_{j,d'}} \right)^2 \end{aligned} \quad (57)$$

is the second-order residual that reduces to the imbalance of *gradients*.² This further implies that near stationarity of the D -gated objective, $\Delta_{j,1,d'}^{(t)}$ also vanishes.

For the scalar-scalar imbalances $\mathcal{I}_{j,d,d'}^{(t+1)}$, we already derived $(\gamma_{j,d'}^{(t+1)})^2$ in Eq. (55), which also symmetrically defines $(\gamma_{j,d}^{(t+1)})^2$. Subtracting both, the first-order terms cancel again, and we obtain

$$\mathcal{I}_{j,d,d'}^{(t+1)} = \left(1 - \frac{4\lambda}{D} \eta \right) \mathcal{I}_{j,d,d'}^{(t)} + \eta^2 ((A_{j,d}^{(t)})^2 - (A_{j,d'}^{(t)})^2) := \left(1 - \frac{4\lambda}{D} \eta \right) \mathcal{I}_{j,d,d'}^{(t)} + \eta^2 \Delta_{j,d,d'}^{(t)}, \quad (58)$$

where $(A_{j,d}^{(t)})^2 := \left(\prod_{d'' \neq d} \gamma_{j,d''}^{(t)} \cdot \omega_j^{(t)\top} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} \gamma_{j,d}^{(t)} \right)^2 = \left(\frac{\partial \mathcal{L}(\omega, \Gamma)}{\partial \gamma_{j,d}} \right)^2$. This shows point i).

Similar to the vector-scalar case, $\Delta_{j,d,d'}^{(t)}$ being a difference of squared partial derivatives implies that the term vanishes at stationarity.

For point ii), we assume balancedness, i.e., $\mathcal{I}_{j,d,d'}^{(t)} = 0$, for any two scalars $d, d' > 1$. The imbalance under (S)GD evolves as $\mathcal{I}_{j,d,d'}^{(t+1)} = \left(1 - \frac{4\lambda}{D} \eta \right) \mathcal{I}_{j,d,d'}^{(t)} + \eta^2 ((A_{j,d}^{(t)})^2 - (A_{j,d'}^{(t)})^2)$. Therefore, we must show that pair-wise balancedness $\mathcal{I}_{j,d,d'}^{(t)} = 0$ (with potentially non-zero factors $\gamma_{j,d}, \gamma_{j,d'}$) implies $\Delta_{j,d,d'}^{(t+1)} = 0$ for $d, d' > 1$. Denoting $(\gamma_j^{\odot \setminus d})^{(t)} := \prod_{d'' \neq d} \gamma_{j,d''}^{(t)}$, we factor the second-order term as

$$\Delta_{j,d,d'}^{(t)} = (A_{j,d}^{(t)})^2 - (A_{j,d'}^{(t)})^2 = (A_{j,d}^{(t)} + A_{j,d'}^{(t)})(A_{j,d}^{(t)} - A_{j,d'}^{(t)}). \quad (59)$$

The sum term is $((\gamma_j^{\odot \setminus d})^{(t)} + (\gamma_j^{\odot \setminus d'})^{(t)}) \cdot \omega_j^{(t)\top} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} (\gamma_{j,d}^{(t)} + \gamma_{j,d'}^{(t)})$, and the difference term $((\gamma_j^{\odot \setminus d})^{(t)} - (\gamma_j^{\odot \setminus d'})^{(t)}) \cdot \omega_j^{(t)\top} \mathbf{g}_j^{(t)} + \frac{2\lambda}{D} (\gamma_{j,d}^{(t)} - \gamma_{j,d'}^{(t)})$. Expanding their product, we obtain:

$$\begin{aligned} \Delta_{j,d,d'}^{(t)} &= \left((\gamma_j^{\odot \setminus d})^{(t)} + (\gamma_j^{\odot \setminus d'})^{(t)} \right) \left((\gamma_j^{\odot \setminus d})^{(t)} - (\gamma_j^{\odot \setminus d'})^{(t)} \right) (\omega_j^{(t)\top} \mathbf{g}_j^{(t)})^2 + \left(\frac{4\lambda^2}{D^2} \right) \left((\gamma_{j,d}^{(t)})^2 - (\gamma_{j,d'}^{(t)})^2 \right) \\ &\quad + \left(\frac{2\lambda}{D} \right) (\gamma_{j,d}^{(t)} + \gamma_{j,d'}^{(t)}) \left((\gamma_j^{\odot \setminus d})^{(t)} - (\gamma_j^{\odot \setminus d'})^{(t)} \right) \omega_j^{(t)\top} \mathbf{g}_j^{(t)} \\ &\quad + \left(\frac{2\lambda}{D} \right) (\gamma_{j,d}^{(t)} - \gamma_{j,d'}^{(t)}) \left((\gamma_j^{\odot \setminus d})^{(t)} + (\gamma_j^{\odot \setminus d'})^{(t)} \right) \omega_j^{(t)\top} \mathbf{g}_j^{(t)}. \end{aligned} \quad (60)$$

²Note that this can be further simplified to $\Delta_{j,1,d'}^{(t)} = ((\gamma_j^{\odot \setminus d'})^{(t)})^2 \cdot \left((\gamma_{j,d'}^{(t)})^2 \|\mathbf{g}_j^{(t)}\|_2^2 - (\omega_j^{(t)\top} \mathbf{g}_j^{(t)})^2 \right) + \frac{4\lambda^2}{D^2} \mathcal{I}_{j,1,d'}^{(t)}$, showing that we can not fully factor $\mathcal{I}_{j,1,d'}^{(t)}$ out of $\Delta_{j,1,d'}^{(t)}$.

The last two mixed terms cancel, which can be seen by summing them and factoring:

$$\begin{aligned} & \left(\frac{2\lambda}{D} \right) \boldsymbol{\omega}_j^{(t)\top} \mathbf{g}_j^{(t)} \cdot \left[\left(\gamma_{j,d}^{(t)} + \gamma_{j,d'}^{(t)} \right) \left((\gamma_j^{\odot \setminus d})^{(t)} - (\gamma_j^{\odot \setminus d'})^{(t)} \right) \right. \\ & \quad \left. + \left(\gamma_{j,d}^{(t)} - \gamma_{j,d'}^{(t)} \right) \left((\gamma_j^{\odot \setminus d})^{(t)} + (\gamma_j^{\odot \setminus d'})^{(t)} \right) \right] \end{aligned} \quad (61)$$

Now, expanding the bracket, all terms but $2\gamma_{j,d}^{(t)}(\gamma_j^{\odot \setminus d})^{(t)} - 2\gamma_{j,d'}^{(t)}(\gamma_j^{\odot \setminus d'})^{(t)} = 2(\gamma_j^{\odot})^{(t)} - 2(\gamma_j^{\odot})^{(t)} = 0$ cancel, so that the remaining expression is

$$\Delta_{j,d,d'}^{(t)} = \left(((\gamma_j^{\odot \setminus d})^{(t)})^2 - ((\gamma_j^{\odot \setminus d'})^{(t)})^2 \right) (\boldsymbol{\omega}_j^{(t)\top} \mathbf{g}_j^{(t)})^2 + \left(\frac{4\lambda^2}{D^2} \right) \mathcal{I}_{j,d,d'}^{(t)}. \quad (62)$$

Finally, using $(\gamma_j^{\odot \setminus d})^{(t)} = \gamma_{j,d'}^{(t)} \cdot (\gamma_j^{\odot \setminus d,d'})^{(t)} := \gamma_{j,d'}^{(t)} \cdot \prod_{d'' \in [D-1] \setminus \{d,d'\}} \gamma_{j,d''}^{(t)}$, we can simplify

$$\begin{aligned} \Delta_{j,d,d'}^{(t)} &= \left(((\gamma_j^{\odot \setminus d})^{(t)})^2 - ((\gamma_j^{\odot \setminus d'})^{(t)})^2 \right) (\boldsymbol{\omega}_j^{(t)\top} \mathbf{g}_j^{(t)})^2 + \left(\frac{4\lambda^2}{D^2} \right) \mathcal{I}_{j,d,d'}^{(t)} \\ &= ((\gamma_{j,d'}^{(t)})^2 - (\gamma_{j,d}^{(t)})^2) ((\gamma_j^{\odot \setminus d,d'})^{(t)})^2 (\boldsymbol{\omega}_j^{(t)\top} \mathbf{g}_j^{(t)})^2 + \left(\frac{4\lambda^2}{D^2} \right) \mathcal{I}_{j,d,d'}^{(t)} \\ &= \mathcal{I}_{j,d,d'}^{(t)} \left[-(\boldsymbol{\omega}_j^{(t)\top} \mathbf{g}_j^{(t)})^2 \cdot ((\gamma_j^{\odot \setminus d,d'})^{(t)})^2 + \frac{4\lambda^2}{D^2} \right] \end{aligned} \quad (63)$$

This shows for scalar gating parameters $d, d' \in [D] \setminus \{1\}$, all terms in $\mathcal{I}_{j,d,d'}^{(t+1)}$ depend multiplicatively on $\mathcal{I}_{j,d,d'}^{(t)}$, and hence $\mathcal{I}_{j,d,d'}^{(t)} = 0$ implies $\mathcal{I}_{j,d,d'}^{(t+1)} = (1 - \frac{4\lambda}{D}) \cdot 0 + \eta^2 \cdot 0 = 0$, proving point ii).

For point iii), note that a balanced zero representation $(\boldsymbol{\omega}_j^{(t)}, \{\gamma_{j,d}^{(t)}\}_{d=1}^{D-1}) = \mathbf{0}$ satisfies $\mathcal{I}_{j,d,d'}^{(t)} = 0$ for all $d, d' \in [D]$. Since all terms in $\Delta_{j,d,d'}^{(t+1)}$ are multiplicative in either $\boldsymbol{\omega}_j^{(t)}$ or $\gamma_{j,d}^{(t)}$, it vanishes. Therefore, $\mathcal{I}_{j,d,d'}^{(t+1)} = (1 - \frac{4\lambda}{D}\eta) \cdot 0 + \eta^2 \cdot 0 = 0$ for $d, d' \neq 1$. Analogously, $\Delta_{j,1,d'}^{(t+1)} = 0$ using the same argument. Hence $\mathcal{I}_{j,1,d'}^{(t+1)} = 0$. Beyond simply enforcing balancedness between both vector and scalar gating factors, inspecting the updates for $\boldsymbol{\omega}_j^{(t)} = \mathbf{0}$ and $\gamma_{j,d}^{(t)} = 0$ in Eq. (53) directly shows that $\boldsymbol{\omega}_j^{(t+1)} = \mathbf{0} - \eta \cdot \mathbf{0} = \mathbf{0}$ and $\gamma_{j,d}^{(t+1)} = 0 - \eta \cdot 0 = 0$. By induction, the iterates of a balanced zero representation in group $j \in [J]$ are thus confined to $\mathbf{w}_j^{(t')} = \mathbf{0} \forall t' > t$. This completes the proof. $\square$

C Experimental details

In all our experiments, we evaluate D -Gating only for $D \in \{2, 3, 4\}$, resulting in induced $L_{2,1}$, $L_{2,2/3}$, and $L_{2,0.5}$ regularization, following broadly established ranges for $L_{p,q}$ group regularization [21]. Although an ever deeper D -Gating parametrization better approximates a group L_0 penalty, the increase in non-convexity potentially causes numerical instabilities without necessarily improving performance beyond 4-Gating (cf. Appendix E.3). Nevertheless, in the majority of our experiments, there is a marked distinction between convex group penalization ($D = 2$) and the non-convex extension ($D > 2$). Throughout all experiments, $\boldsymbol{\omega}$ is initialized using a standard initialization scheme as for $\mathbf{w}$, while the gating parameters $\gamma_{j,d}$ are initialized with unity, i.e., $\boldsymbol{\Gamma}^{(0)} = \mathbf{1}_J \mathbf{1}_{D-1}^\top$, see Algorithm 1.

C.1 Details on group lasso simulation

Simulation details For our group sparse linear regression set-up, we simulate $n = 200$ training (and 2000 test) samples with $p = 200$ features grouped into 40 groups of 5 features each. 7 feature groups contain informative (non-zero) weights and 33 groups contain noise features only. We draw both the feature values and the (informative) weights from $\mathcal{N}(0, 1)$ and likewise add standard Gaussian additive noise to the simulated predictor to obtain the targets.

Table 2: Summary of datasets used in feature selection experiments.

Dataset	Training Samples	Test Samples	Classes	Input Features
ISOLET [8]	6,328	1,559	26	617
COIL20 [42]	1,152	288	20	400
ACTIVITY [1]	4,252	1,492	6	561
MNIST [32]	60,000	10,000	10	784
F-MNIST [62]	60,000	10,000	10	784
MADELON [16]	2,080	520	2	500

Training details We train the D -Gated models and direct $L_{2,1}$ penalization using full-batch gradient descent for 1500 iterations using a cosine learning rate schedule with initial learning rate 5×10^{-2} and 0.9 momentum. We consider parameter groups with norm smaller than 10^{-6} to be zero. For the specialized optimization routine to solve the Group Lasso efficiently, we use the accelerated proximal gradient method proposed by [48]. The *oracle* method shown in Fig. 4 is the efficient ordinary least squares estimator computed only on the truly non-zero coefficients. Within each simulation, all regularized methods are run on a grid of 30 λ values spaced logarithmically between 10^{-5} and 15. The procedure is repeated over 10 simulations and the average results reported in Fig. 4.

C.2 Details on feature selection experiments

Architecture As a backbone for all methods (D -Gating, LassoNet, HSIC), we use a fully-connected LeNet-300-100 architecture [33] with two hidden layers (300 and 100 neurons) and ReLU activation functions and Kaiming Normal initialization. To obtain feature selection for non-linear models using our approach, we group the first-layer weights of the network by input feature and apply D -Gating. We train the network on all classification tasks and methods using SGD for 100 epochs with a batch size of 256 and cosine schedule with an initial learning rate of 0.1. All models are initialized using a Kaiming Normal scheme, with the scaling factors for D -Gating being initialized with unity (cf. Algorithm 1). Moreover, we consider parameter groups with 2-norm smaller than `float32.mach.eps` $\approx 1.19 \times 10^{-7}$ to be 0. Table 2 describes the basic properties of the used datasets. For datasets which do not explicitly provide a test set, we randomly assign 20% of the samples to the test set. Note that, to ensure a fair comparison to D -Gating, LassoNet [35] is implemented without an additional debiasing step that retrains the model only on the selected features.

C.3 Details on filter sparsity experiments

Architectures VGG-16 for CIFAR-10 and SVHN consists of 13 convolutional layers and 3 fully-connected layers [49]. The convolutional layers are organized into five groups: two layers with 64 filters, two with 128 filters, 3 layers with 256 filters, 3 layers with 512 filters, and another 3 layers with 512 filters. After each group, we apply a 2×2 max-pooling operation. All filters are 3×3 . We add batch normalization immediately before every ReLU activation. To reduce parameters compared to the ImageNet version, we follow [65] and replace the two dense layers before the output with a single fully-connected layer of 512 neurons.

ResNet-18 is an 18-layer deep residual network introduced by [18]. For small-image datasets, we follow, e.g., [52], by replacing the initial 7×7 convolution and the following max-pooling layer with a single 3×3 . The model begins with that convolution, then proceeds through four stages of basic residual blocks, each stage consisting of two blocks, with filter counts of [64, 128, 256, 512]. Finally, global average pooling reduces the feature maps before the fully connected output layer. In all methods we only regularize the convolutional but not the fully-connected layers, either through D -Gating or by directly regularizing the filter weights using an $L_{2,1}$ penalty [60, 47, 36] or the batch normalization scaling variables as in network slimming [38].

Data We conduct experiments on three standard image classification benchmark tasks described in Table 3. We use the train/test split provided by the datasets. Additionally, we use standard pre-processing and data augmentation techniques, namely normalizing the input images and employing horizontal flips, width or height shifts up to 12.5%, and up to 15° rotations.

Table 3: Summary of datasets used in image classification experiments.

Dataset	Training Samples	Test Samples	Classes	Input Features
CIFAR-10 [29]	50,000	10,000	10	3,072 ($32 \times 32 \times 3$)
SVHN [43]	73,257	26,032	10	3,072 ($32 \times 32 \times 3$)
CIFAR-100 [29]	50,000	10,000	100	3,072 ($32 \times 32 \times 3$)

Training Our training hyperparameter settings generally follow broadly established configurations for the respective tasks [49, 18, 65]. Because established learning rate settings were found to perform suboptimally, we scanned the interval $[10^{-3}, 1]$ using a small regularization strength λ for each method and dataset separately. Table 4 lists the hyperparameters for the filter sparsity experiments.

Table 4: Training hyperparameters for different architectures and image classification datasets. The learning rates correspond to D -Gating with $D = \{2, 3, 4\}$ (set) and the comparison methods (second value). The comparison methods use standard Kaiming initialization.

Architecture	Dataset	Epochs	Batch size	Optim.	Mom.	Init.	LR	Schedule
VGG-16	CIFAR-10	200	128	SGD	0.9	Kaiming Normal	$\{0.3, 0.4, 0.4\}, 0.1$	Cosine
	SVHN	200	128	SGD	0.9	Kaiming Normal	$\{2e-3, 3e-3, 3e-3\}, 2e-3$	Cosine
ResNet-18	CIFAR-100	200	128	SGD	0.9	Kaiming Normal	$\{0.2, 0.3, 0.4\}, 0.2$	Cosine

In D -Gating, the sparsity is controlled by the regularization strength λ . To obtain the accuracy-sparsity tradeoffs, we therefore train models for each $D \in \{2, 3, 4\}$ along a grid of λ values that spans the whole sparsity range. For the comparison methods, direct optimization of the sparsity-inducing $L_{2,1}$ and L_1 penalties does not result in sparse filters whose norm is below `float32.mach.eps` $\approx 1.19 \times 10^{-7}$. To achieve sparsity and construct the tradeoff curves, we also train these models along a grid of λ values but subsequently post-hoc prune each of these models along a sequence of pre-defined pruning ratios $\{0.1, 0.2, \dots, 0.9, 0.95, 0.98, 0.99\}$. The tradeoff curves are then constructed by taking the best performance over all λ values separately at each pruning ratio. This optimally reflects different regularization requirements at different sparsity ranges. In contrast to the original proposals of [60, 36, 38], we implement both direct $L_{2,1}$ regularization with post-hoc filter pruning and network slimming (L_1 regularization of the batch normalization scaling parameters) as a one-shot procedure without retraining after sparsification to ensure a fair comparison to D -Gating, which can be seen as merely another optimization vehicle to optimize non-smooth structured sparsity penalties, making them comparable outside of a sophisticated pruning pipeline. Moreover, like direct optimization of these penalties, D -Gating can be arbitrarily combined with other pruning and retraining schemes.

C.4 Details on language modeling experiments

We choose a variant of NanoGPT trained on the character-level TinyShakespeare dataset comprising 5,458,199 tokens and a vocabulary size of 91 different symbols. Deviating from the usual performance evaluation metrics for language models, the character-level granularity also allows for reporting the validation accuracy instead of only the perplexity. We set aside 10% of the training data for validation purposes.




Figure 9: Schematic illustration of D -Gating overparametrization to achieve structured head sparsity in a multi-head attention layer. **Left:** original attention layer. **Right:** D -Gated value matrices sparsify the whole attention head. Red nodes indicate the location of the additional gating parameters.

Our NanoGPT implementation contains 10.8 mio. parameters and has 6 layers, 12 attention heads per multi-head attention layer, block size of 256, an embedding dimension of 384 and dropout

with parameter 0.2 between the layers. We train all models on a grid of logarithmically spaced λ values between 10^{-4} and 10^2 for 3000 iterations using a batch size of 32 and Adam [24] with a cosine schedule and initial learning rate 10^{-3} . Since direct optimization of the $L_{2,1}$ penalty does not yield sparse solutions, defined as having a 2-norm smaller than 10^{-5} , we further prune these models using structured global magnitude pruning to sparsity ratios $\{0.05, 0.1, 0.2, \dots, 0.9, 0.95\}$, and report the best performance over all λ values for a given sparsity ratio. Note that the D -Gated models incorporate no additional pruning.

Fig. 9 illustrates how we implement structured attention head penalization using D -Gating. Note that we only gate the value matrices of each head, which suffices for the whole attention head to be removed when sparsified.

C.5 Details on Neural Oblivious Decision Ensembles

We optimize the squared loss on the Wine quality prediction task ($n = 4,898; p = 11$) [6] using SGD with Nesterov momentum 0.9 and a cosine decay learning rate schedule with initial learning rate 0.4. The models are trained for 400 epochs without early stopping using a batch size of 1024. Fig. 10 describes the application of D -Gating in NODEs for $D = 2$.




Figure 10: Schematic illustration of D -Gating overparametrization to achieve structured tree sparsity in a NODE model. **Left:** original NODE structure. **Right:** NODE with 2-Gated tree structures (multiplied by a shared factor).

D Computational Environment

All experiments were conducted either on a single NVIDIA RTX A6000 or A4000 GPU with 48GB and 16GB of memory, respectively. Smaller experiments, e.g., for toy models, were carried out on standard CPU workstations. The total single-GPU runtime for all experiments is estimated to be around 500 hours.

E Further experiments and ablation studies

E.1 Additional plots for filter sparsity in image classification

Besides the accuracy-sparsity tradeoff, the reduction of FLOPs through is of special interest for structured sparsification methods. Besides Table 1, we show the full tradeoff curves for accuracy and theoretical speed-up, measured as the ratio of FLOPs of the original and sparse models, in Fig. 11.

Additionally, Fig. 12 shows the regularization paths for D -Gating with $D \in \{2, 3, 4\}$ for all three image classification datasets we experimented on. Complementing the conclusions from the tradeoff curves in the main text, both plots corroborate our findings of superior tradeoffs for $D = 3, 4$ over $D = 2$, as well as increased sparsity at the same regularization parameter λ .

E.2 Overheads

Parameter overhead To quantify the parameter overhead incurred by applying D -Gating in our experiments, we contrast the number of additional parameters with the models' parameters. The results are shown in Table 5, highlighting that larger architectures typically have fewer regularized groups (attention heads, convolutional filters) with large group sizes of thousands of parameters,




Figure 11: Test accuracy vs. theoretical speed-up (measured as the ratio of FLOPs) of D -Gated convolutional neural networks and comparison methods.




Figure 12: Effect of regularization λ on filter sparsity for $D \in \{2, 3, 4\}$.

suggesting that our approach scales better for larger than for small architectures by introducing relatively fewer additional parameters.

Table 5: Parameter count overhead for D -Gating. The linear model is the one from Fig. 4, the input-sparse LeNet-300-100, the one from Fig. 5, VGG-16 from Fig. 6, and NanoGPT from Fig. 7. The overhead rate in the worst case ($D = 4$ and the largest input dimension), is given by $\frac{\text{\#new params}}{\text{\#old params}} - 1$.

Model	Vanilla Params	Additional Params	Overhead rate ($D = 4$)	Details
Lin. Mod.	200	$40 \cdot (D-1)$	6.0×10^{-1}	40 groups of 5 features
LeNet-300-100	$\approx 266,000$	$\text{input_dim} \cdot (D-1)$	8.8×10^{-3}	$\text{input_dim} \in [400, 784]$
VGG-16	≈ 15 mio.	$4,224 \cdot (D-1)$	8.4×10^{-4}	Sum of conv. filters in 13 layers
NanoGPT	≈ 10.8 mio.	$72 \cdot (D-1)$	2.0×10^{-5}	6 layers layers, 12 att. heads

Runtime and memory overhead Additionally, we measure the wall-clock runtime overhead of D -Gating for $D \in \{2, 3, 4\}$ against the vanilla models. Similarly, we record the peak GPU memory utilization during training.

The runtime results are shown in Fig. 13, showing slight to moderate increases between 0% and 30% for smaller batch sizes, with the difference becoming increasingly irrelevant for commonly used larger batch sizes as well as larger architectures.

Similarly, Fig. 14 shows the memory overhead over vanilla models, indicating indiscernible utilization for the vast majority of settings. The behavior of the fully-connected LeNet-300-100 exhibits high variance and less clear patterns compared to the larger models, potentially due to its small size. Generally, part of the difference in runtime and memory overhead is likely due to implementation efficiency reasons and not necessarily an inherent difference.

E.3 Depth and instability ablation studies

To guide the selection of the gating depth D and demonstrate that the full benefits of non-convex $L_{2,q}$, $q < 1$, regularization are usually attained at $q = 2/3$ or $q = 0.5$, corresponding to $D = 3, 4$ for D -Gating [21], we further conduct depth ablation studies for the input-sparse LeNet-300-100 from the experiment in Fig. 5 on ISOLET, extend the results for the filter-sparse VGG-16 on CIFAR-10 (Fig. 6), and analyze performance and instability for the group sparse regression task in Fig. 4.




Figure 13: Wall-clock runtime for $D \in \{2, 3, 4\}$ and vanilla model. Means and std. over ten runs are shown.




Figure 14: Peak memory use for $D \in \{2, 3, 4\}$ and vanilla model. Means and std. over ten runs are shown.

Fig. 15 shows the results for LeNet-300-100 and VGG-16 for increasing depth up to $D = 20$ and $D = 8$, respectively. Importantly, we observe no measurable improvement in the performance-sparsity tradeoff beyond $D = 4$, while the instability increases slightly.

For the group sparse linear regression task shown in Fig. 16, the left panel illustrates the numerical instability caused by too large D , evidenced by the jagged regularization paths for $D > 6$. The tradeoff curves in the right panel reveal that these instabilities degrade the performance at lower sparsity values for large D .

Taken together, the values of $q \in \{2/3, 0.5\}$, i.e., $D \in \{3, 4\}$, recommended in the literature are supported by our ablation studies, where deeper gating approaches show at best equivalent performance while increasing numerical instability.

E.4 Imbalance decay and loss convergence for SGD and Adam

The convergence of the D -Gating objective $\mathcal{L}(\omega, \Gamma)$ to the non-smooth regularized objective $\mathcal{L}_{\mathbf{w}}(\mathbf{w})$, as shown in Fig. 3, is directly related to the convergence of the regularizers of both objectives, as well as the balancedness of the gating representations. To see this, note that

$$\begin{aligned}
 \mathcal{L}(\omega, \Gamma) - \mathcal{L}_{\mathbf{w}}(\mathbf{w}) &= \mathcal{L}_0(\omega \triangleright \gamma^\odot) + \lambda(\|\omega\|_2^2 + \|\Gamma\|_F^2)/D - \mathcal{L}_0(\mathbf{w}) - \lambda\|\mathbf{w}\|_{2,2/D}^{2/D} \\
 &= \lambda((\|\omega\|_2^2 + \|\Gamma\|_F^2)/D - \|\mathbf{w}\|_{2,2/D}^{2/D}) \\
 &= \lambda(\mathcal{R}(\omega, \Gamma) - \mathcal{R}_{\mathbf{w}}(\mathbf{w})) \\
 &= \lambda\mathcal{M}(\omega, \Gamma),
 \end{aligned} \tag{64}$$

Hence, the losses converge if the regularizers converge, which, by the AM-GM inequality, happens if and only if the gating representations are balanced. This relationship is verified in Fig. 17, showing the same experiment as in Fig. 3, but plotting the convergence of regularizers instead of losses.

Another interesting question is to empirically investigate the evolution of misalignment for a common optimizer, for which the gradient flow and descent results Lemma 3 and Lemma 4 do not apply, i.e., Adam [24]. Fig. 18 shows the evolution of misalignment for the same setting as the previous Fig. 17, but optimized with Adam instead of SGD. Results show that the D -Gating objective also converges to the sparsity-inducing objective $\mathcal{L}_{\mathbf{w}}(\mathbf{w})$ faster than SGD overall, albeit not at a neat exponential rate. Moreover, the convergence speeds do not seem to depend as strongly on D as for SGD.

E.5 Subnetwork selection

In an additional application, we investigate whether D -Gating can correctly select whole data modalities in a simple late-fusion multimodal network [12]. In these architectures, the modalities are processed independently in separate subnetworks before their latent representations are fused in




Figure 15: Effect of increasing the gating depth D beyond $D = 4$ for the input-sparse LeNet-300-100 on ISOLET as well as the filter-sparse VGG-16 trained on CIFAR-10.




Figure 16: Effect of increasing the gating depth D beyond $D = 4$ on the group-sparse linear regression task. The left panel shows the regularization paths, and the right panel the performance-sparsity tradeoffs. The setup is the same as in Fig. 4.

the penultimate late-fusion layer. Fig. 19 visualizes such an architecture as well as our D -Gating proposal for differentiable selection of modalities. In the D -Gated variant, the final-layer weights are grouped by modality and subsequently gated, effectively selecting out whole modality subnetworks. To control the contribution of the modalities, we simulate a semi-synthetic dataset combining both tabular and image information additively in a predictor. To be precise, we simulate targets as

$$Y = (1 - \alpha) \cdot \eta^{\text{tabular}} + \alpha \cdot \eta^{\text{image}} + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, 1), \quad \alpha \in \{0, 0.5, 1\}.$$

The image predictor is based on the MNIST dataset of handwritten images and is simply the (demeaned) face value of the digit contained in the image. The tabular predictor is given by

$$\eta^{\text{tabular}} = \sum_{k=1}^6 f_k(x_k), \quad x_k \sim \mathcal{U}[-1, 1],$$

where the $f_k(\cdot)$ are fixed smooth non-linear shape functions. Both predictors are constructed to be similarly distributed. We simulate $n = 2000$ training and test observations from the above data-generating process. In the late-fusion architecture, we use a small VGG-style subnetwork with 4 convolutional layers for the MNIST images and a single neural additive subnetwork to model the smooth non-linear effects $f_k(\cdot)$ of the tabular covariates x_k . The models are optimized for 300 epochs using SGD with 0.8 momentum, a batch size of 32, and a learning rate of 5×10^{-3} .

Fig. 20 shows the results for $D \in \{2, 3, 4\}$ and $\alpha \in \{0, 0.5, 1\}$. As expected, using the one standard error rule to select the regularization strength λ , D -Gating can consistently select out the unimportant modalities, with the separation of important and unimportant modalities becoming clearer as D increases. Notably, for equal modality contributions (middle row), neither the validation-optimal λ_{opt} nor $\lambda_{1\text{se}}$ selects out any of the modalities.




Figure 17: Convergence of regularizers, i.e., evolution of misalignment, instead of loss convergence for SGD trained on a neuron-sparse LeNet-300-100. Setting is the same as for Fig. 3.




Figure 18: Convergence of regularizers with Adam instead of SGD trained on a neuron-sparse LeNet-300-100. Setting is the same as for Fig. 3.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: Yes, we clearly state our claims in the abstract and contribution subsection. The claims are then further substantiated in the main sections using formal proofs validated using experiments.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: We include a specific subsection on the limitations of our proposed method in our discussion section and also discuss limitations of our theoretical analysis and experiments in the main sections.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors




Figure 19: Schematic illustration of modality selection in a late-fusion multimodal network using D -Gating overparametrization. **Left:** original late-fusion architecture. **Right:** architecture including D -Gating of the late-fusion layer, where the red nodes indicate the location of the gating factors.




Figure 20: Regularization path and weight norms in multi-modal late-fusion network for different modality contributions (**rows**) and D -Gating depth $D \in \{2, 3, 4\}$ (**columns**). The solid vertical bar indicates the validation-optimal λ_{opt} and the dashed vertical bar the λ_{1se} chosen by the 1se rule, i.e., the sparsest model that performs within one standard error of the lowest validation RMSE (10 splits).

should reflect on how these assumptions might be violated in practice and what the implications would be.

- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: We provide complete proofs with the full set of assumptions in Appendix B, with the proven statements containing the full assumptions.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: We provide a detailed description of our experimental set-up for all experiments, including hyperparameter choices and their justification in Appendix C. All our experiments are done on well-known and openly accessible benchmark datasets. We also formalize our method in Algorithm 1 for reproducibility.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All used datasets are publicly available and the source code for the experiments and baselines will be submitted in the supplementary materials. We aim to collect the code in a publicly accessible Git repository as soon as possible.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: We provide a detailed description of our experimental setup, used hyperparameters, and their justification in Appendix C. Finer details can be found in the full code submitted in the supplementary material.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: We provide error bars and their description (e.g., standard deviations, standard errors, or confidence intervals) for all experiments where multiple runs are not prohibitively expensive (e.g., NanoGPT).

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We specify our computational environment in Appendix D and provide an estimate of the total runtime of the experiments.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification:

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification:

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification:

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: We reference all datasets and disclose the version of the used packages in the dependencies of the submitted code.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification:

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification:

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification:

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or originality of the research, declaration is not required.

Answer: [NA]

Justification:

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.